

a02

May 3, 2024

```
[1]: # Edited by: Kok Teng Ng (1936360), Minjeong Lee (1978925)
      # IE 678 Deep Learning, University of Mannheim
      # Author: Rainer Gemulla
```

1 Deep Learning: Assignment 2

```
[2]: %matplotlib ipynb

      # Define imports & defaults
      import numpy as np
      import torch
      import torch.nn as nn
      from torch.utils.data import DataLoader, Dataset

      # import helper functions
      import os
      import sys

      sys.path.append(os.getcwd())
      from a02helper import *

      DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
      torch.manual_seed(0) # ensure reproducibility
      np.random.seed(0)

      MAX_SEQ_LEN = 200
      BATCH_SIZE = 32
```

Warning: Cannot change to a different GUI toolkit: notebook. Using ipynb instead.

1.1 Task 1: Datasets

```
[3]: import string
      from torchtext.vocab import vocab
      from torchtext.data import get_tokenizer
```

```

class ReviewsDataset(Dataset):
    def __init__(
        self,
        reviews_file="data/reviews_small.txt",
        labels_file="data/labels_small.txt",
        use_vocab=False,
    ):
        """
        A dataset of movie reviews and their labels.

        Args:
            reviews_file: the reviews file
            labels_file: the labels file
            use_vocab: if True, yield reviews in a numerical representation
        """
        # Load data from filesystem
        with open(reviews_file) as f:
            raw_reviews = f.readlines()

        with open(labels_file) as f:
            raw_labels = f.readlines()

        # Preprocessing and store (in memory)
        self._reviews = self._preprocess_reviews(raw_reviews)
        self._labels = self._preprocess_labels(raw_labels)

        # Build vocabulary
        self.vocab = None
        if use_vocab:
            from torchtext.vocab import build_vocab_from_iterator
            self.vocab = build_vocab_from_iterator(
                self._reviews, specials=["<pad>"] # will get token id 0
            )

    def __len__(self):
        """Returns the length of the dataset."""
        return len(self._reviews)

    def __getitem__(self, idx):
        """
        Returns a tuple of a preprocessed review and the corresponding label.
        ↪ If the
        vocabulary is enabled, returns a numerical representation of the review.

        Args:
            idx: a single index
        Returns: a (review, label) tuple
        """

```

```

        """
        review = self._reviews[idx]
        label = self._labels[idx]
        if self.vocab is not None:
            #review_ids = self.vocab([token for token in review])
            review_ids = [self.vocab[token] for token in review if token in
↪self.vocab]
            return review_ids, label
        else:
            return review, label

def _preprocess_reviews(self, raw_reviews):
    """
    Applies two kinds of preprocessing:

    (i) Apply the "basic_english" tokenizer from the torchtext library to
    transform every review into a list of normalized tokens (cf.
    https://pytorch.org/text/stable/data_utils.html#get-tokenizer).

    (ii) Remove punctuation (cf.
    https://docs.python.org/3/library/string.html#string.punctuation).

    Returns: list of tokenized reviews
    """
    tokenizer = get_tokenizer("basic_english")
    punctuation = set(string.punctuation)

    processed_reviews = []
    for review in raw_reviews:
        processed_review = tokenizer(review.strip().lower())
        processed_review = [tok for tok in processed_review if tok not in
↪punctuation]
        processed_reviews.append(processed_review)
    return processed_reviews

def _preprocess_labels(self, raw_labels):
    """
    Transform raw labels into integers, where 1="positive" and 0 otherwise.

    Returns: list of labels
    """
    # Hint: You can remove leading and trailing whitespace from the raw
↪labels using
    # the strip() method.
    preprocessed_labels_list = []
    for label in raw_labels:
        label = label.strip()

```

```

        if label == "positive":
            preprocessed_labels_list.append(1)
        else:
            preprocessed_labels_list.append(0)
    return preprocessed_labels_list

```

[4]: *# Test your code (without vocabulary).*

```

dataset = ReviewsDataset()
print(dataset[0])

```

Should yield:

(['bromwell', 'high', 'is', 'a', 'cartoon', 'comedy', ...], 1)

```

(['bromwell', 'high', 'is', 'a', 'cartoon', 'comedy', 'it', 'ran', 'at', 'the',
'same', 'time', 'as', 'some', 'other', 'programs', 'about', 'school', 'life',
'such', 'as', 'teachers', 'my', 'years', 'in', 'the', 'teaching', 'profession',
'lead', 'me', 'to', 'believe', 'that', 'bromwell', 'high', 's', 'satire', 'is',
'much', 'closer', 'to', 'reality', 'than', 'is', 'teachers', 'the', 'scramble',
'to', 'survive', 'financially', 'the', 'insightful', 'students', 'who', 'can',
'see', 'right', 'through', 'their', 'pathetic', 'teachers', 'pomp', 'the',
'pettiness', 'of', 'the', 'whole', 'situation', 'all', 'remind', 'me', 'of',
'the', 'schools', 'i', 'knew', 'and', 'their', 'students', 'when', 'i', 'saw',
'the', 'episode', 'in', 'which', 'a', 'student', 'repeatedly', 'tried', 'to',
'burn', 'down', 'the', 'school', 'i', 'immediately', 'recalled', 'at', 'high',
'a', 'classic', 'line', 'inspector', 'i', 'm', 'here', 'to', 'sack', 'one',
'of', 'your', 'teachers', 'student', 'welcome', 'to', 'bromwell', 'high', 'i',
'expect', 'that', 'many', 'adults', 'of', 'my', 'age', 'think', 'that',
'bromwell', 'high', 'is', 'far', 'fetched', 'what', 'a', 'pity', 'that', 'it',
'isn', 't'], 1)

```

[5]: *# Test your code (with vocabulary).*

```

dataset = ReviewsDataset(use_vocab=True)
print(dataset[0])

```

Should yield:

([10661, 307, 6, 3, 1177, 202, 8, ...], 1)

```

([10661, 307, 6, 3, 1177, 202, 8, 2248, 33, 1, 168, 56, 15, 49, 85, 8902, 43,
422, 122, 140, 15, 3234, 59, 144, 9, 1, 5504, 6267, 454, 72, 5, 260, 12, 10661,
307, 13, 2060, 6, 73, 2780, 5, 692, 76, 6, 3234, 1, 29527, 5, 1730, 7117, 1,
6161, 1726, 36, 52, 68, 212, 143, 63, 1409, 3234, 17974, 1, 28056, 4, 1, 221,
758, 31, 2748, 72, 4, 1, 6311, 10, 731, 2, 63, 1726, 54, 10, 208, 1, 321, 9, 64,
3, 1601, 4042, 743, 5, 2853, 187, 1, 422, 10, 1254, 10116, 33, 307, 3, 380, 322,
6162, 10, 135, 136, 5, 10172, 30, 4, 134, 3234, 1601, 2545, 5, 10661, 307, 10,
529, 12, 113, 1841, 4, 59, 676, 103, 12, 10661, 307, 6, 227, 4163, 48, 3, 2201,
12, 8, 231, 21], 1)

```

1.2 Task 2: Data Loaders

```
[6]: # Split dataset into training, validation, and test subsets
dataset = ReviewsDataset(use_vocab=True)
train_set, val_set, test_set = torch.utils.data.random_split(
    dataset, [0.8, 0.1, 0.1], generator=torch.Generator().manual_seed(123)
)
```

1.2.1 Task 2a

```
[7]: # Example usage of a data loader
dataloader = DataLoader(
    val_set, # a dataset
    1, # desired batch size
    False, # whether to randomly shuffle the dataset
    num_workers = 0, # number of workers that construct batches in parallel
)
```

```
[8]: # Let's print the first batch
batch = next(iter(dataloader))
print(batch)
# [[tensor([11]), tensor([6]), tensor([1]), ..., tensor([0])]
```

```

[tensor([11]), tensor([6]), tensor([1]), tensor([1037]), tensor([6578]),
tensor([4]), tensor([10]), tensor([89]), tensor([120]), tensor([48]),
tensor([163]), tensor([47]), tensor([6]), tensor([27]), tensor([342]),
tensor([4]), tensor([2228]), tensor([140]), tensor([3]), tensor([186]),
tensor([1466]), tensor([1]), tensor([771]), tensor([26]), tensor([78]),
tensor([1459]), tensor([200]), tensor([1101]), tensor([1]), tensor([66]),
tensor([26]), tensor([78]), tensor([5199]), tensor([5]), tensor([2288]),
tensor([1]), tensor([7861]), tensor([6591]), tensor([2]), tensor([83]),
tensor([2446]), tensor([25]), tensor([1]), tensor([182]), tensor([2756]),
tensor([2520]), tensor([34]), tensor([1]), tensor([145]), tensor([8]),
tensor([13]), tensor([39]), tensor([290]), tensor([25]), tensor([252]),
tensor([14480]), tensor([32]), tensor([52]), tensor([497]), tensor([9]),
tensor([223]), tensor([1]), tensor([3254]), tensor([25]), tensor([937]),
tensor([2]), tensor([153]), tensor([568]), tensor([5]), tensor([91]),
tensor([2]), tensor([30]), tensor([388]), tensor([1110]), tensor([17]),
tensor([80]), tensor([62]), tensor([1]), tensor([119]), tensor([255]),
tensor([14]), tensor([34]), tensor([2632]), tensor([1539]), tensor([133]),
tensor([28]), tensor([6]), tensor([31]), tensor([56]), tensor([33]),
tensor([27]), tensor([119]), tensor([413]), tensor([1]), tensor([230]),
tensor([4]), tensor([1036]), tensor([17]), tensor([3]), tensor([738]),
tensor([552]), tensor([1305]), tensor([11189]), tensor([14923]), tensor([6]),
tensor([88]), tensor([203]), tensor([3]), tensor([50]), tensor([283]),
tensor([19]), tensor([1]), tensor([27078]), tensor([44]), tensor([683]),
tensor([8]), tensor([0])]
```

```
[9]: print(len(val_set[0][0]))
      print(len(val_set[1][0]))
      print(len(val_set[2][0]))
```

116

158

131

```
[10]: # explanation: for the Question 2a

# The current code in cell [7] treats each review (a row of input sequence
↳data) as an individual batch. This becomes problematic because reviews can
↳have different lengths. By default,
# PyTorch's DataLoader( ) expects all elements in a batch to be the same size.
↳When the batch size increases and we mix reviews of varying lengths, the
↳DataLoader( ) encounters errors.

# To overcome this challenge, we can define a maximum sequence length for the
↳input data. This essentially sets a cap on the length of each review. If a
↳review exceeds the limit, we can
# truncate it by keeping only the first words up to the maximum length.
↳Conversely, shorter reviews can be padded with a special token (like
↳"<PAD>", in our case we are using 0 instead of
# "<PAD>" based on the provided code setting within the ReviewDataset function)
↳to ensure all reviews within a batch have the same length without causing
↳error from DataLoader. This
# padding technique effectively increase the length of shorter reviews to match
↳the maximum, allowing the DataLoader to function smoothly.

# By defining a maximum sequence length and handling variable-length reviews
↳through truncation and padding, we ensure the DataLoader can efficiently
↳process batches of reviews during
# training.
```

1.2.2 Task 2b

```
[11]: def review_collate_fn(raw_batch):
      """Prepare batches of reviews from a review dataset.

      Args:
          raw_batch: collection of (review, label)-tuples from a ReviewDataset

      Returns: a tuple (review x token id tensor, label tensor) of sizes
          batch_size*MAX_SEQ_LEN and batch_size, respectively.

      """
      reviews, labels = zip(*raw_batch)
```

```

padded_reviews = torch.zeros(len(reviews), MAX_SEQ_LEN, dtype=torch.long)

for i, review in enumerate(reviews):
    if len(review) <= MAX_SEQ_LEN:
        padded_reviews[i, :len(review)] = torch.tensor(review)
    else:
        padded_reviews[i, :] = torch.tensor(review[:MAX_SEQ_LEN])
label_tensor = torch.tensor(labels, dtype=torch.long)
return padded_reviews, label_tensor

```

```

[12]: # Test your function
review_collate_fn([[1, 2, 3], 1), (torch.arange(MAX_SEQ_LEN * 2) + 1, 0)])

# Should yield:
# (tensor([[ 1,  2,  3,  0,  0, ..., 0 ],
#          [ 1,  2,  3,  4,  5, ..., 200 ]]),
#  tensor([1, 0]))

```

/var/folders/2h/v65x911156j8d_r4wnykkqrw0000gn/T/ipykernel_3491/3521082536.py:18
: UserWarning: To copy construct from a tensor, it is recommended to use
sourceTensor.clone().detach() or
sourceTensor.clone().detach().requires_grad_(True), rather than
torch.tensor(sourceTensor).

```

    padded_reviews[i, :] = torch.tensor(review[:MAX_SEQ_LEN])

```

```

[12]: (tensor([[ 1,  2,  3,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0],
                [ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14,
                15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28,
                29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42,
                43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56,
                57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70,
                71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84,
                85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98,
                99, 100, 101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 112,

```

```

113, 114, 115, 116, 117, 118, 119, 120, 121, 122, 123, 124, 125, 126,
127, 128, 129, 130, 131, 132, 133, 134, 135, 136, 137, 138, 139, 140,
141, 142, 143, 144, 145, 146, 147, 148, 149, 150, 151, 152, 153, 154,
155, 156, 157, 158, 159, 160, 161, 162, 163, 164, 165, 166, 167, 168,
169, 170, 171, 172, 173, 174, 175, 176, 177, 178, 179, 180, 181, 182,
183, 184, 185, 186, 187, 188, 189, 190, 191, 192, 193, 194, 195, 196,
197, 198, 199, 200]]),
tensor([1, 0]))

```

```

[13]: # Create the data loaders (with shuffling for training data -> randomization)
train_loader = DataLoader(train_set, BATCH_SIZE, True,
    ↳collate_fn=review_collate_fn)
val_loader = DataLoader(val_set, BATCH_SIZE, False,
    ↳collate_fn=review_collate_fn)
test_loader = DataLoader(test_set, BATCH_SIZE, False,
    ↳collate_fn=review_collate_fn)

```

```

[14]: # Let's print the first batch
batch = next(iter(val_loader))
print(batch)

# (tensor([[ 11,    6,    1, ...,  0,    0,    0],
#          [ 11,   170,  2220, ...,  0,    0,    0],
#          [ 48,    3, 30376, ...,  0,    0,    0],
#          ...,
#          [ 176,   56,   10, ...,  0,    0,    0],
#          [ 239,   534,  1404, ...,  44,   120,    1],
#          [2954, 15576,    6, ..., 2678,   65,    1]]),
# tensor([0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0,
# ↳1,
#          0, 0, 0, 1, 1, 1, 0, 0]))

```

```

(tensor([[ 11,    6,    1, ...,  0,    0,    0],
          [ 11,   170,  2220, ...,  0,    0,    0],
          [ 48,    3, 30376, ...,  0,    0,    0],
          ...,
          [ 176,   56,   10, ...,  0,    0,    0],
          [ 239,   534,  1404, ...,  44,   120,    1],
          [2954, 15576,    6, ..., 2678,   65,    1]]), tensor([0, 1, 0, 1,
0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1,
0, 0, 0, 1, 1, 1, 0, 0]))

```


1.3 Task 3: Recurrent Neural Networks

```
[15]: class SimpleLSTM(nn.Module):
    def __init__(
        self, vocab_size, embedding_dim, hidden_dim, num_layers=1,
        cell_dropout=0.0
    ):
        """
        Initializes the model by setting up the layers.

        Args:
            vocab_size: number of unique words in the reviews
            embedding_dim: size of the embeddings
            hidden_dim: dimension of the LSTM output
            num_layers: number of LSTM layers
            cell_dropout: dropout applied between the LSTM layers
                        (provide to LSTM constructor as dropout argument)
        """
        super().__init__()
        self.num_layers = num_layers
        self.hidden_dim = hidden_dim
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.cell_dropout = cell_dropout
        self.lstm = nn.LSTM(embedding_dim, hidden_dim, batch_first = True,
        num_layers = num_layers, dropout = cell_dropout)
        self.fc = nn.Linear(hidden_dim, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        """
        Performs a forward pass of the model on some input and hidden state.

        Parameters
        -----
        x: batch as a (batch_size, sequence_length) tensor

        Returns
        -----
        Probability of positive class and the last output of the LSTM.
        """
        hidden = self.init_hidden(len(x))
        embeddings = self.embedding(x)
        outputs, (hidden_state, cell_state) = self.lstm(embeddings, hidden)
        thought_vector = hidden_state[-1, :, :]
        logits = self.fc(thought_vector)
        predictions = self.sigmoid(logits)
        return predictions, thought_vector
```

```

def init_hidden(self, batch_size):
    """Initialize hidden states.

    Returns a tuple of two num_layers x batch_size x hidden_dim tensors
    (one for initial cell states, one for initial hidden states) consisting of all
    zeros.
    """
    # Note: ensure that the returned tensors are located on DEVICE.
    hidden = (torch.zeros(self.num_layers, batch_size, self.hidden_dim,
        device = self.embedding.weight.device),
        torch.zeros(self.num_layers, batch_size, self.hidden_dim,
        device = self.embedding.weight.device))
    return hidden

# Test constructor
model = SimpleLSTM(50, 10, 32, 2, 0.1).to(DEVICE)
print(model)

# Should give:
# SimpleLSTM(
#   (embedding): Embedding(50, 10)
#   (lstm): LSTM(10, 32, num_layers=2, batch_first=True, dropout=0.1)
#   (fc): Linear(in_features=32, out_features=1, bias=True)
#   (sigmoid): Sigmoid()
# )

```

```

SimpleLSTM(
  (embedding): Embedding(50, 10)
  (lstm): LSTM(10, 32, num_layers=2, batch_first=True, dropout=0.1)
  (fc): Linear(in_features=32, out_features=1, bias=True)
  (sigmoid): Sigmoid()
)

```

```

[16]: # Test forward pass
model = SimpleLSTM(50, 10, 32, 2).to(DEVICE)
dummy_data = torch.arange(30, dtype=torch.int, device=DEVICE).reshape(3, 10)

# fix model parameters
for key in model.state_dict():
    model.state_dict()[key][:] = 0.1
probs, states = model(dummy_data)
print(probs)
print(states)

```

```

# tensor([[0.9643],
#         [0.9643],
#         [0.9643]], device='cuda:0 or cpu', grad_fn=<SigmoidBackward0>)
# tensor([[0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
#         [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
#         [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.
↪9985,
#         0.9985, 0.9985, 0.9985, 0.9985, 0.9985], device='cuda:0 or cpu',
#         grad_fn=<SliceBackward0>)

```

```

tensor([[0.9643],
        [0.9643],
        [0.9643]], grad_fn=<SigmoidBackward0>)
tensor([[0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
        [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985],
        [0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985, 0.9985,
        0.9985, 0.9985, 0.9985, 0.9985, 0.9985]], grad_fn=<SliceBackward0>)

```

1.3.1 Task 3d

```
[17]: @torch.no_grad() # Disables autograd for this function
def reviews_eval(
    model, eval_loader, label="val", loss_fn=torch.nn.functional.
    ↪binary_cross_entropy
):
    model.to(DEVICE)
    model.eval() # sets model to evaluation mode (e.g., relevant for dropout)

    total_correct = total_loss = 0
    for reviews, labels in eval_loader:
        reviews, labels = reviews.to(DEVICE), labels.to(DEVICE)

        # Forward pass: Compute the model's output, reshape it to a vector, and
        # then run the provided loss function.
        prediction, _ = model(reviews)
        prediction = prediction.view(-1)
        loss = loss_fn(prediction, labels.float())

        # Eval stats: Add loss to total_loss and number of correct predictions
    ↪to
        # total_correct.
        total_loss = total_loss + loss.item()
        prediction_class = torch.round(prediction)
        correct = (prediction_class == labels).sum().item()
        total_correct = total_correct + correct

    print(
        f"          "
        f"{label} accuracy: {total_correct / len(eval_loader.dataset):.4f}\t"
        f"{label} loss: {total_loss / len(eval_loader):.4f}"
    )
```

```
[18]: # Test your implementation
model = SimpleLSTM(len(dataset.vocab), 10, 10, 1, 0).to(DEVICE)
reviews_eval(model, val_loader)

# Should yield with different but similar numbers:
#          val accuracy: 0.5100          val loss: 0.6928
```

val accuracy: 0.5375 val loss: 0.6925

1.3.2 Task 3e

```
[19]: def reviews_train(
    model,
    train_loader,
    val_loader,
    lr=0.01,
    epochs=3,
    max_norm=5,
    loss_fn=torch.nn.functional.binary_cross_entropy,
):
    """
    Train a network on the review data

    Args:
        model: Initialized model
        train_loader: Dataloader for the training data
        val_loader: Dataloader for the validation data
        lr: learning rate
        epochs: number of epochs
        max_norm: max norm of gradients for gradient clipping
        loss_fn: Loss function
    """
    # Send the model's parameters to your accelerator (cuda or cpu)
    model = model.to(DEVICE)

    # Define optimizer for the parameters which require gradients (cf. Task 5)
    optimizer = torch.optim.Adam(
        [param for param in model.parameters() if param.requires_grad], lr=lr
    )

    # Let's go
    for epoch in range(epochs):
        total_correct = total_loss = 0
        for reviews, labels in train_loader:
            model.train()

            # Send batch to your accelerator
            reviews, labels = reviews.to(DEVICE), labels.to(DEVICE)

            # Forward pass: Compute the model's output, reshape it to a vector,
            ↪ and then
            # run the provided loss function.
            prediction, _ = model(reviews)
            prediction = prediction.view(-1)
            loss = loss_fn(prediction, labels.float())
```

```

        # Backward pass:
        # (i) Compute the gradients wrt. the loss
        # (ii) Clip the gradients using
        #      https://pytorch.org/docs/stable/generated/torch.nn.utils.
        ↪ clip_grad_norm_.html to max_norm
        # (iii) Run the optimizer
        # (iv) Clear all accumulated gradients
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm = ↪
        ↪ max_norm)
        optimizer.step()
        optimizer.zero_grad()

        # Compute epoch statistics:
        # (i) Add the loss of this batch to the total_loss
        # (ii) Add the number of correct predictions (max prob) to ↪
        ↪ total_correct
        total_loss = total_loss + loss.item()
        prediction_class = torch.round(prediction)
        correct = (prediction_class == labels).sum().item()
        total_correct = total_correct + correct

    print(
        f"Epoch {epoch + 1:2}\t"
        f"train accuracy: {total_correct / len(train_loader.dataset):.4f}\t"
        f"train loss: {total_loss / len(train_loader):.4f}"
    )

    # now validate
    reviews_eval(model, val_loader, loss_fn=loss_fn)

```

```

[20]: # Test your implementation
model = SimpleLSTM(len(dataset.vocab), 10, 10, 1).to(DEVICE)
reviews_train(model, train_loader, val_loader, epochs=5)

# Should yield something like (note: numbers have high variance over runs):
# Epoch 1      train accuracy: 0.4994      train loss: 0.6953
#              val accuracy: 0.4875      val loss: 0.6922
# Epoch 2      train accuracy: 0.5319      train loss: 0.6885
#              val accuracy: 0.5275      val loss: 0.6861
# Epoch 3      train accuracy: 0.6059      train loss: 0.6443
#              val accuracy: 0.5400      val loss: 0.6902
# Epoch 4      train accuracy: 0.6863      train loss: 0.5438
#              val accuracy: 0.5925      val loss: 0.7453
# Epoch 5      train accuracy: 0.8334      train loss: 0.3875
#              val accuracy: 0.7300      val loss: 0.6310

```

Epoch 1	train accuracy: 0.5131	train loss: 0.6940
	val accuracy: 0.4800	val loss: 0.6956
Epoch 2	train accuracy: 0.5503	train loss: 0.6854
	val accuracy: 0.5650	val loss: 0.6881
Epoch 3	train accuracy: 0.5953	train loss: 0.6339
	val accuracy: 0.5475	val loss: 0.7067
Epoch 4	train accuracy: 0.6653	train loss: 0.5385
	val accuracy: 0.5225	val loss: 0.7984
Epoch 5	train accuracy: 0.7106	train loss: 0.4592
	val accuracy: 0.5975	val loss: 0.8588

1.4 Task 4: Pre-trained Embeddings & Visualization

1.4.1 Task 4b

```
[21]: # Load Glove embeddings into a plain embedding layer.
vocab = dataset.vocab
glove_embeddings = nn.Embedding(len(vocab), 100, device=DEVICE)
reviews_load_embeddings(glove_embeddings, vocab.get_stoi())
```


Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

```
[22]: # Print one embedding
glove_embeddings(torch.tensor(vocab["movie"], device=DEVICE))
```

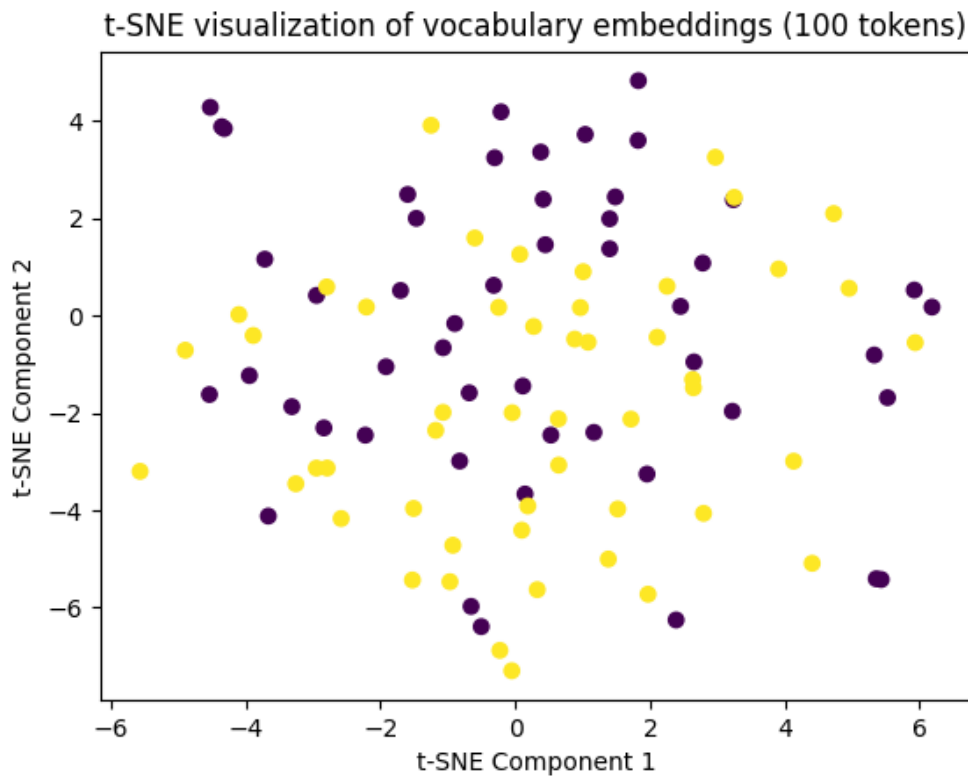
```
[22]: tensor([ 0.3825,  0.1482,  0.6060, -0.5153,  0.4399,  0.0611, -0.6272, -0.0254,
          0.1643, -0.2210,  0.1442, -0.3721, -0.2168, -0.0890,  0.0979,  0.6561,
          0.6446,  0.4770,  0.8385,  1.6486,  0.8892, -0.1181, -0.0125, -0.5208,
          0.7785,  0.4872, -0.0150, -0.1413, -0.3475, -0.2959,  0.1028,  0.5719,
         -0.0456,  0.0264,  0.5382,  0.3226,  0.4079, -0.0436, -0.1460, -0.4835,
          0.3204,  0.5509, -0.7626,  0.4327,  0.6175, -0.3650, -0.6060, -0.7962,
          0.3929, -0.2367, -0.3472, -0.6120,  0.5475,  0.9481,  0.2094, -2.7771,
         -0.6022,  0.8495,  1.2549,  0.0179, -0.0419,  2.1147, -0.0266, -0.2810,
          0.6812, -0.1417,  0.9925,  0.4988, -0.6754,  0.6417,  0.4230, -0.2791,
          0.0634,  0.6891, -0.3618,  0.0537, -0.1681,  0.1942, -0.4707, -0.1480,
         -0.5899, -0.2797,  0.1679,  0.1057, -1.7601,  0.0088, -0.8333, -0.5836,
         -0.3708, -0.5659,  0.2070,  0.0713,  0.0556, -0.2976, -0.0727, -0.2560,
          0.4269,  0.0589,  0.0911,  0.4728], grad_fn=<EmbeddingBackward0>)
```

```
[23]: # Plot embeddings of first 100 words using t-SNE
nextplot()
_ = tsne_vocab(glove_embeddings, torch.arange(100), vocab)
```

t-SNE visualization of vocabulary embeddings (100 tokens)



```
[24]: # You can also specify colors and/or drop the item labels
nextplot()
_ = tsne_vocab(glove_embeddings, torch.arange(100), colors=[0] * 50 + [1] * 50)
```





```
[25]: import umap

embeddings = glove_embeddings.weight.detach().cpu().numpy()
indices = np.random.choice(range(len(vocab)), 100, replace=False)
selected_embeddings = embeddings[indices]
selected_words = [vocab.get_itos()[i] for i in indices]
reducer = umap.UMAP(n_neighbors=5, n_components=2, metric='cosine')
umap_embeddings = reducer.fit_transform(selected_embeddings)
plt.figure(figsize=(6, 6))

plt.scatter(umap_embeddings[:, 0], umap_embeddings[:, 1], s=1)
plt.title('2D UMAP projection of Selected GloVe Embeddings')
for i, word in enumerate(selected_words):
    plt.annotate(word, (umap_embeddings[i, 0], umap_embeddings[i, 1]),
        ↪fontsize=4)
plt.show()
```

2D UMAP projection of Selected GloVe Embeddings



1.4.2 Task 4c

```
[26]: # hyperparameter settings for rest of task 4
vocab_size = len(dataset.vocab)
embedding_dim = 100
hidden_dim = 100
num_layers = 2
n_epochs = 10
cell_dropout = 0.0
```

```
[27]: # train a plain model
model = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
                  cell_dropout).to(
```

```

    DEVICE
)
reviews_train(model, train_loader, val_loader, epochs=n_epochs)

# Should reach a (train) accuracy of >0.9. If not, rerun.

```

```


Epoch 1      train accuracy: 0.4928  train loss: 0.6999
              val accuracy: 0.4825    val loss: 0.6957
Epoch 2      train accuracy: 0.4928  train loss: 0.6938
              val accuracy: 0.4825    val loss: 0.6934
Epoch 3      train accuracy: 0.5106  train loss: 0.6929
              val accuracy: 0.4725    val loss: 0.6939
Epoch 4      train accuracy: 0.5316  train loss: 0.6897
              val accuracy: 0.5000    val loss: 0.6998
Epoch 5      train accuracy: 0.6009  train loss: 0.6364
              val accuracy: 0.5200    val loss: 0.7398
Epoch 6      train accuracy: 0.6500  train loss: 0.5464
              val accuracy: 0.4950    val loss: 0.8098
Epoch 7      train accuracy: 0.7472  train loss: 0.4570
              val accuracy: 0.6175    val loss: 0.8547
Epoch 8      train accuracy: 0.8853  train loss: 0.2865
              val accuracy: 0.6675    val loss: 0.8698
Epoch 9      train accuracy: 0.9500  train loss: 0.1435
              val accuracy: 0.7125    val loss: 0.8566
Epoch 10     train accuracy: 0.9812  train loss: 0.0560
              val accuracy: 0.7350    val loss: 0.9660

```

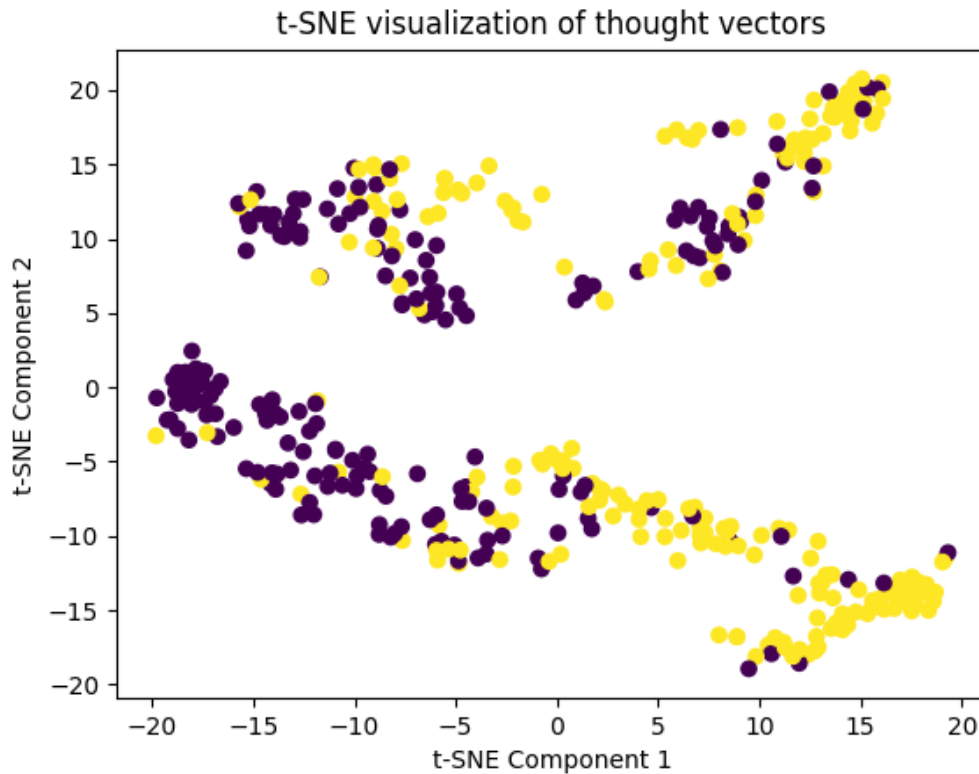
```

[28]: # Plot t-SNE embeddings of the thought vectors for training data
      # point color = label
      nextplot()
      _ = tsne_thought(model, train_loader, DEVICE)

```



```
[29]: # Plot t-SNE embeddings of the thought vectors for validation data
nextplot()
_ = tsne_thought(model, val_loader, DEVICE)
```



1.4.3 Task 4d

```
[30]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
    ↪ cell_dropout)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.4888	train loss: 0.6950
	val accuracy: 0.4850	val loss: 0.6932
Epoch 2	train accuracy: 0.5153	train loss: 0.6905
	val accuracy: 0.5200	val loss: 0.6928
Epoch 3	train accuracy: 0.5141	train loss: 0.6928
	val accuracy: 0.4850	val loss: 0.6932
Epoch 4	train accuracy: 0.5425	train loss: 0.6826
	val accuracy: 0.4950	val loss: 0.6943
Epoch 5	train accuracy: 0.6425	train loss: 0.5942
	val accuracy: 0.6900	val loss: 0.6207
Epoch 6	train accuracy: 0.8213	train loss: 0.3973

	val accuracy: 0.7550	val loss: 0.5638
Epoch 7	train accuracy: 0.9031	train loss: 0.2660
	val accuracy: 0.7825	val loss: 0.5483
Epoch 8	train accuracy: 0.9337	train loss: 0.1904
	val accuracy: 0.7950	val loss: 0.6733
Epoch 9	train accuracy: 0.9547	train loss: 0.1314
	val accuracy: 0.7650	val loss: 0.7809
Epoch 10	train accuracy: 0.9716	train loss: 0.0880
	val accuracy: 0.8000	val loss: 0.7238

1.4.4 Task 4e

```
[31]: # Initialize the model with pre-trained embeddings without finetuning, then
      ↪ train
model_p = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
      ↪ cell_dropout)
reviews_load_embeddings(model_p.embedding, vocab.get_stoi())
model_p.embedding.weight.requires_grad = False
reviews_train(model_p, train_loader, val_loader, epochs=n_epochs)
```

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.4928	train loss: 0.6951
	val accuracy: 0.5125	val loss: 0.6931
Epoch 2	train accuracy: 0.5141	train loss: 0.6929
	val accuracy: 0.5200	val loss: 0.6919
Epoch 3	train accuracy: 0.5162	train loss: 0.6896
	val accuracy: 0.4875	val loss: 0.8306
Epoch 4	train accuracy: 0.5262	train loss: 0.6957
	val accuracy: 0.4900	val loss: 0.6936
Epoch 5	train accuracy: 0.5509	train loss: 0.6783
	val accuracy: 0.5100	val loss: 0.6943
Epoch 6	train accuracy: 0.5766	train loss: 0.6511
	val accuracy: 0.5025	val loss: 0.7334
Epoch 7	train accuracy: 0.6075	train loss: 0.6163
	val accuracy: 0.5825	val loss: 0.7070
Epoch 8	train accuracy: 0.6972	train loss: 0.5623
	val accuracy: 0.7125	val loss: 0.6833
Epoch 9	train accuracy: 0.7309	train loss: 0.5567
	val accuracy: 0.6875	val loss: 0.6195
Epoch 10	train accuracy: 0.7503	train loss: 0.5176
	val accuracy: 0.6875	val loss: 0.6246

1.5 Task 5: Exploration

Task 5a

```
[32]: class SimpleLSTM(nn.Module):
      def __init__(
```

```

        self, vocab_size, embedding_dim, hidden_dim, num_layers=1, cell_type =
↪ "lstm", cell_dropout=0.0, bidirectional=False
    ):
        super().__init__()
        self.num_layers = num_layers
        self.hidden_dim = hidden_dim
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.cell_dropout = cell_dropout
        self.cell_type = cell_type
        self.bidirectional = bidirectional
        if self.cell_type == "lstm":
            self.rnn = torch.nn.LSTM(embedding_dim, hidden_dim, num_layers =
↪ num_layers, batch_first = True, dropout = cell_dropout, bidirectional=self.
↪ bidirectional)
        elif self.cell_type == "gru":
            self.rnn = torch.nn.GRU(embedding_dim, hidden_dim, num_layers =
↪ num_layers, batch_first = True, dropout = cell_dropout, bidirectional=self.
↪ bidirectional)
        elif self.cell_type == "elman":
            self.rnn = torch.nn.RNN(embedding_dim, hidden_dim, num_layers =
↪ num_layers, batch_first = True, dropout = cell_dropout, bidirectional=self.
↪ bidirectional)
        self.fc = nn.Linear(hidden_dim * 2 if self.bidirectional else
↪ hidden_dim, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        hidden = self.init_hidden(len(x))
        embeddings = self.embedding(x)
        outputs, hidden_state = self.rnn(embeddings, hidden)
        if self.bidirectional:
            if self.cell_type == "lstm":
                hidden_state = torch.cat((hidden_state[0][-2,:,:],
↪ hidden_state[0][-1,:,:]), dim = 1)
            else:
                hidden_state = torch.cat((hidden_state[-2,:,:],
↪ hidden_state[-1,:,:]), dim = 1)
            else:
                if self.cell_type == "lstm":
                    hidden_state = hidden_state[0][-1,:,:]
                else:
                    hidden_state = hidden_state[-1,:,:]
        thought_vector = hidden_state
        logits = self.fc(thought_vector)
        predictions = self.sigmoid(logits)
        return predictions, thought_vector

```

```

def init_hidden(self, batch_size):
    num_directions = 2 if self.bidirectional else 1
    if self.cell_type == "lstm":
        hidden = (torch.zeros(self.num_layers * num_directions, batch_size,
↪self.hidden_dim, device = self.embedding.weight.device),
        torch.zeros(self.num_layers * num_directions, batch_size,
↪self.hidden_dim, device = self.embedding.weight.device))
    else:
        hidden = torch.zeros(self.num_layers * num_directions, batch_size,
↪self.hidden_dim, device = self.embedding.weight.device)
    return hidden

```

dropout = 0.0 and Bidirectional LSTM

```

[33]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
↪"lstm", 0.0, True)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)

```

Initializing embedding layer with pretrained word embeddings...


Initialized 29841/32363 word embeddings

Epoch 1	train accuracy: 0.5153	train loss: 0.6936
	val accuracy: 0.5200	val loss: 0.6912
Epoch 2	train accuracy: 0.6497	train loss: 0.6190
	val accuracy: 0.6600	val loss: 0.6246
Epoch 3	train accuracy: 0.8825	train loss: 0.2839
	val accuracy: 0.8225	val loss: 0.4654
Epoch 4	train accuracy: 0.9825	train loss: 0.0585
	val accuracy: 0.8725	val loss: 0.4033
Epoch 5	train accuracy: 0.9975	train loss: 0.0158
	val accuracy: 0.8125	val loss: 0.7838
Epoch 6	train accuracy: 0.9991	train loss: 0.0060
	val accuracy: 0.8850	val loss: 0.5905
Epoch 7	train accuracy: 0.9997	train loss: 0.0026
	val accuracy: 0.8625	val loss: 0.6895
Epoch 8	train accuracy: 0.9997	train loss: 0.0022
	val accuracy: 0.8650	val loss: 0.7565
Epoch 9	train accuracy: 0.9997	train loss: 0.0028
	val accuracy: 0.8725	val loss: 0.6960
Epoch 10	train accuracy: 0.9997	train loss: 0.0023
	val accuracy: 0.8700	val loss: 0.7447


```

[34]: nextplot()
_ = tsne_thought(model_pf, train_loader, DEVICE)

```

```
[35]: nextplot()  
_ = tsne_thought(model_pf, val_loader, DEVICE)
```



dropout = 0.0 and Unidirectional LSTM

```
[36]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers, ↵
↵ "lstm", 0.0, False)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```


Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1	train accuracy: 0.5059	train loss: 0.6947
	val accuracy: 0.4825	val loss: 0.6934
Epoch 2	train accuracy: 0.5587	train loss: 0.6769
	val accuracy: 0.5200	val loss: 0.7043
Epoch 3	train accuracy: 0.7847	train loss: 0.4528
	val accuracy: 0.7250	val loss: 0.5789
Epoch 4	train accuracy: 0.9453	train loss: 0.1758
	val accuracy: 0.8150	val loss: 0.5023
Epoch 5	train accuracy: 0.9825	train loss: 0.0616
	val accuracy: 0.7900	val loss: 0.6953
Epoch 6	train accuracy: 0.9944	train loss: 0.0234
	val accuracy: 0.8100	val loss: 0.7986

Epoch 7	train accuracy: 0.9950	train loss: 0.0162
	val accuracy: 0.8025	val loss: 0.8413
Epoch 8	train accuracy: 0.9984	train loss: 0.0065
	val accuracy: 0.8275	val loss: 0.8040
Epoch 9	train accuracy: 0.9994	train loss: 0.0037
	val accuracy: 0.8125	val loss: 1.0207
Epoch 10	train accuracy: 0.9997	train loss: 0.0028
	val accuracy: 0.8150	val loss: 1.1340

```
[37]: nextplot()
      _ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[38]: nextplot()
      _ = tsne_thought(model_pf, val_loader, DEVICE)
```



dropout = 0.8 and Bidirectional LSTM

```
[39]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers, ↵
↵ "lstm", 0.8, True)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```


Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1	train accuracy: 0.5094	train loss: 0.6977
	val accuracy: 0.5300	val loss: 0.6923
Epoch 2	train accuracy: 0.5559	train loss: 0.6848
	val accuracy: 0.5950	val loss: 0.6508
Epoch 3	train accuracy: 0.7547	train loss: 0.5102
	val accuracy: 0.8450	val loss: 0.4315
Epoch 4	train accuracy: 0.9403	train loss: 0.1704
	val accuracy: 0.8200	val loss: 0.6849
Epoch 5	train accuracy: 0.9803	train loss: 0.0611
	val accuracy: 0.8450	val loss: 0.4435
Epoch 6	train accuracy: 0.9938	train loss: 0.0248
	val accuracy: 0.8500	val loss: 0.7188

Epoch 7	train accuracy: 0.9959	train loss: 0.0162
	val accuracy: 0.8625	val loss: 0.7143
Epoch 8	train accuracy: 0.9931	train loss: 0.0218
	val accuracy: 0.8775	val loss: 0.5299
Epoch 9	train accuracy: 0.9966	train loss: 0.0112
	val accuracy: 0.8600	val loss: 0.5964
Epoch 10	train accuracy: 0.9981	train loss: 0.0081
	val accuracy: 0.8675	val loss: 0.5266

```
[40]: nextplot()
      _ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[41]: nextplot()
      _ = tsne_thought(model_pf, val_loader, DEVICE)
```



dropout = 0.8 and Unidirectional LSTM

```
[42]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
↳ "lstm", 0.8, False)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```


Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1	train accuracy: 0.5044	train loss: 0.6961
	val accuracy: 0.4825	val loss: 0.6964
Epoch 2	train accuracy: 0.5034	train loss: 0.6961
	val accuracy: 0.4875	val loss: 0.6954
Epoch 3	train accuracy: 0.5122	train loss: 0.6952
	val accuracy: 0.5350	val loss: 0.6921
Epoch 4	train accuracy: 0.5238	train loss: 0.6924
	val accuracy: 0.4900	val loss: 0.7053
Epoch 5	train accuracy: 0.5781	train loss: 0.6610
	val accuracy: 0.5075	val loss: 0.7048
Epoch 6	train accuracy: 0.6547	train loss: 0.5851
	val accuracy: 0.6300	val loss: 0.7208

Epoch 7	train accuracy: 0.7581	train loss: 0.4761
	val accuracy: 0.6900	val loss: 0.6710
Epoch 8	train accuracy: 0.8544	train loss: 0.3498
	val accuracy: 0.7350	val loss: 0.6609
Epoch 9	train accuracy: 0.8862	train loss: 0.2898
	val accuracy: 0.6250	val loss: 1.0919
Epoch 10	train accuracy: 0.9087	train loss: 0.2460
	val accuracy: 0.7675	val loss: 0.6516

```
[43]: nextplot()
      _ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[44]: nextplot()
      _ = tsne_thought(model_pf, val_loader, DEVICE)
```



dropout = 0.0 and Unidirectional GRU cell

```
[45]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers, "gru",
                      ↪0.0, False)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```


Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1	train accuracy: 0.5766	train loss: 0.6873
	val accuracy: 0.4800	val loss: 0.8398
Epoch 2	train accuracy: 0.8137	train loss: 0.3855
	val accuracy: 0.8700	val loss: 0.3228
Epoch 3	train accuracy: 0.9819	train loss: 0.0645
	val accuracy: 0.8700	val loss: 0.3085
Epoch 4	train accuracy: 0.9962	train loss: 0.0185
	val accuracy: 0.8725	val loss: 0.4523
Epoch 5	train accuracy: 0.9997	train loss: 0.0039
	val accuracy: 0.8700	val loss: 0.5667
Epoch 6	train accuracy: 0.9997	train loss: 0.0017
	val accuracy: 0.8725	val loss: 0.6787


```
Epoch 7      train accuracy: 0.9994  train loss: 0.0021
            val accuracy: 0.8725     val loss: 0.6353
Epoch 8      train accuracy: 1.0000  train loss: 0.0003
            val accuracy: 0.8700     val loss: 0.6769
Epoch 9      train accuracy: 1.0000  train loss: 0.0002
            val accuracy: 0.8725     val loss: 0.7060
Epoch 10     train accuracy: 1.0000  train loss: 0.0002
            val accuracy: 0.8725     val loss: 0.7278
```

```
[46]: nextplot()
      _ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[47]: nextplot()
      _ = tsne_thought(model_pf, val_loader, DEVICE)
```



dropout = 0.8 and Bidirectional GRU cell

```
[48]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers, "gru",
                      ↪0.8, True)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```


Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1	train accuracy: 0.6166	train loss: 0.6434
	val accuracy: 0.8375	val loss: 0.3994
Epoch 2	train accuracy: 0.9216	train loss: 0.2040
	val accuracy: 0.8700	val loss: 0.3392
Epoch 3	train accuracy: 0.9928	train loss: 0.0248
	val accuracy: 0.8825	val loss: 0.4144
Epoch 4	train accuracy: 0.9994	train loss: 0.0017
	val accuracy: 0.8875	val loss: 0.4964
Epoch 5	train accuracy: 0.9994	train loss: 0.0027
	val accuracy: 0.8500	val loss: 0.8054
Epoch 6	train accuracy: 0.9997	train loss: 0.0017
	val accuracy: 0.8700	val loss: 0.5759

Epoch 7	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8675	val loss: 0.6222
Epoch 8	train accuracy: 1.0000	train loss: 0.0001
	val accuracy: 0.8700	val loss: 0.6641
Epoch 9	train accuracy: 1.0000	train loss: 0.0000
	val accuracy: 0.8650	val loss: 0.7094
Epoch 10	train accuracy: 1.0000	train loss: 0.0000
	val accuracy: 0.8650	val loss: 0.7323

```
[49]: nextplot()
      _ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[50]: nextplot()
      _ = tsne_thought(model_pf, val_loader, DEVICE)
```



dropout = 0.0 and Unidirectional Elman cell

```
[51]: # Initialize the model with pre-trained embeddings with finetuning, then train
model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
    ↪ "elman", 0.0, False)
reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```


Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


Epoch 1	train accuracy: 0.4994	train loss: 0.7154
	val accuracy: 0.4900	val loss: 0.6986
Epoch 2	train accuracy: 0.5106	train loss: 0.7100
	val accuracy: 0.4900	val loss: 0.7114
Epoch 3	train accuracy: 0.5328	train loss: 0.6935
	val accuracy: 0.4850	val loss: 0.7275
Epoch 4	train accuracy: 0.5909	train loss: 0.6388
	val accuracy: 0.4850	val loss: 0.7406
Epoch 5	train accuracy: 0.6534	train loss: 0.5554
	val accuracy: 0.5100	val loss: 0.9516
Epoch 6	train accuracy: 0.6656	train loss: 0.5044
	val accuracy: 0.4700	val loss: 0.8663

Epoch 7	train accuracy: 0.6947	train loss: 0.4791
	val accuracy: 0.4825	val loss: 0.8753
Epoch 8	train accuracy: 0.6984	train loss: 0.4638
	val accuracy: 0.5025	val loss: 0.9274
Epoch 9	train accuracy: 0.7066	train loss: 0.4414
	val accuracy: 0.5025	val loss: 1.2211
Epoch 10	train accuracy: 0.6934	train loss: 0.4360
	val accuracy: 0.4800	val loss: 1.2877

```
[52]: nextplot()
      _ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[53]: nextplot()
      _ = tsne_thought(model_pf, val_loader, DEVICE)
```



Task 5b

```
[54]: @torch.no_grad()
def reviews_eval(
    model, eval_loader, label="val", loss_fn=torch.nn.functional.
    ↪ binary_cross_entropy
):
    model.to(DEVICE)
    model.eval()

    total_correct = total_loss = 0
    for reviews, labels in eval_loader:
        reviews, labels = reviews.to(DEVICE), labels.to(DEVICE)

        prediction, _ = model(reviews)
        prediction = prediction.view(-1)
        loss = loss_fn(prediction, labels.float())

        total_loss = total_loss + loss.item()
        prediction_class = torch.round(prediction)
        correct = (prediction_class == labels).sum().item()
```

```

        total_correct = total_correct + correct

    val_loss = total_loss / len(eval_loader)
    val_accuracy = total_correct / len(eval_loader.dataset)

    print(
        f"
        {label} accuracy: {val_accuracy:.4f}\t"
        f"{label} loss: {val_loss:.4f}"
    )

    return val_loss

```

```

[55]: def reviews_train(
    model,
    train_loader,
    val_loader,
    lr=0.01,
    epochs=3,
    max_norm=5,
    loss_fn=torch.nn.functional.binary_cross_entropy,
    patience = 3,
):
    model = model.to(DEVICE)

    optimizer = torch.optim.Adam(
        [param for param in model.parameters() if param.requires_grad], lr=lr
    )

    best_val_loss = float("inf")
    patience_counter = 0
    for epoch in range(epochs):
        total_correct = total_loss = 0
        for reviews, labels in train_loader:
            model.train()

            reviews, labels = reviews.to(DEVICE), labels.to(DEVICE)

            prediction, _ = model(reviews)
            prediction = prediction.view(-1)
            loss = loss_fn(prediction, labels.float())

            loss.backward()
            torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm =
↪max_norm)
            optimizer.step()
            optimizer.zero_grad()

```

```

        total_loss = total_loss + loss.item()
        prediction_class = torch.round(prediction)
        correct = (prediction_class == labels).sum().item()
        total_correct = total_correct + correct

    val_loss = reviews_eval(model, val_loader, loss_fn=loss_fn)

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        torch.save(model.state_dict(), f"model_best_val_loss.pt")
        patience_counter = 0
    else:
        patience_counter += 1
    print(
        f"Epoch {epoch + 1:2}\t"
        f"train accuracy: {total_correct / len(train_loader.dataset):.4f}\t"
        f"train loss: {total_loss / len(train_loader):.4f}"
    )
    if patience_counter >= patience:
        print(f"Early stopping triggered after {patience} epochs with no_
↪validation loss improvement.")
        break
    return best_val_loss

```

```

[56]: vocab_size = len(dataset.vocab)
      embedding_dim = 100
      n_epochs = 10
      cell_type = 'gru'
      bidirectional = True
      num_layers = 2

```

```

[57]: import optuna

def objective(trial):
    hidden_dim = trial.suggest_int('hidden_dim', 64, 128)
    cell_dropout = trial.suggest_float('cell_dropout', 0.0, 0.5)
    model = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
↪cell_type, cell_dropout, bidirectional)
    reviews_load_embeddings(model.embedding, vocab.get_stoi())
    best_val_loss = reviews_train(model, train_loader, val_loader,
↪epochs=n_epochs)
    return best_val_loss

study = optuna.create_study(direction='minimize')
study.optimize(objective, n_trials=20)

```

[I 2024-05-03 18:32:07,226] A new study created in memory with name: no-name-b769d2f6-87ce-4d57-8ff1-c4a9f0868e28

Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings

	val accuracy: 0.8350	val loss: 0.4329
Epoch 1	train accuracy: 0.6653	train loss: 0.6229
	val accuracy: 0.8575	val loss: 0.3728
Epoch 2	train accuracy: 0.9378	train loss: 0.1713
	val accuracy: 0.8725	val loss: 0.4074
Epoch 3	train accuracy: 0.9941	train loss: 0.0209
	val accuracy: 0.8775	val loss: 0.5673
Epoch 4	train accuracy: 1.0000	train loss: 0.0010

[I 2024-05-03 18:34:38,980] Trial 0 finished with value: 0.37276275570576006 and parameters: {'hidden_dim': 126, 'cell_dropout': 0.16898837483042878}. Best is trial 0 with value: 0.37276275570576006.

	val accuracy: 0.8775	val loss: 0.5998
Epoch 5	train accuracy: 1.0000	train loss: 0.0001

Early stopping triggered after 3 epochs with no validation loss improvement.

Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings

	val accuracy: 0.8175	val loss: 0.4541
Epoch 1	train accuracy: 0.6328	train loss: 0.6255
	val accuracy: 0.8575	val loss: 0.3690
Epoch 2	train accuracy: 0.9278	train loss: 0.1860
	val accuracy: 0.8525	val loss: 0.5111
Epoch 3	train accuracy: 0.9953	train loss: 0.0152
	val accuracy: 0.8700	val loss: 0.6198
Epoch 4	train accuracy: 1.0000	train loss: 0.0007

[I 2024-05-03 18:37:12,413] Trial 1 finished with value: 0.368980552141483 and parameters: {'hidden_dim': 111, 'cell_dropout': 0.357598769745876}. Best is trial 1 with value: 0.368980552141483.

	val accuracy: 0.8625	val loss: 0.6881
Epoch 5	train accuracy: 1.0000	train loss: 0.0002

Early stopping triggered after 3 epochs with no validation loss improvement.

Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings

	val accuracy: 0.7700	val loss: 0.4559
Epoch 1	train accuracy: 0.6175	train loss: 0.6447
	val accuracy: 0.8800	val loss: 0.2979
Epoch 2	train accuracy: 0.9209	train loss: 0.1980
	val accuracy: 0.8925	val loss: 0.5393
Epoch 3	train accuracy: 0.9972	train loss: 0.0104
	val accuracy: 0.9025	val loss: 0.4824
Epoch 4	train accuracy: 0.9994	train loss: 0.0030

[I 2024-05-03 18:39:49,157] Trial 2 finished with value: 0.29790392288794887 and parameters: {'hidden_dim': 106, 'cell_dropout': 0.3334983015769939}. Best is trial 2 with value: 0.29790392288794887.

```

        val accuracy: 0.9050          val loss: 0.5177
Epoch 5      train accuracy: 1.0000  train loss: 0.0001
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
        val accuracy: 0.8725          val loss: 0.3068
Epoch 1      train accuracy: 0.6653  train loss: 0.5779
        val accuracy: 0.8775          val loss: 0.2906
Epoch 2      train accuracy: 0.9516  train loss: 0.1229
        val accuracy: 0.8425          val loss: 0.6517
Epoch 3      train accuracy: 0.9981  train loss: 0.0070
        val accuracy: 0.8675          val loss: 0.6400
Epoch 4      train accuracy: 0.9997  train loss: 0.0012

[I 2024-05-03 18:41:39,431] Trial 3 finished with value: 0.2906067268206523 and
parameters: {'hidden_dim': 81, 'cell_dropout': 0.31094170398776455}. Best is
trial 3 with value: 0.2906067268206523.

```

```

        val accuracy: 0.8625          val loss: 0.7101
Epoch 5      train accuracy: 1.0000  train loss: 0.0001
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
        val accuracy: 0.8425          val loss: 0.3641
Epoch 1      train accuracy: 0.6728  train loss: 0.5829
        val accuracy: 0.8475          val loss: 0.4024
Epoch 2      train accuracy: 0.9516  train loss: 0.1363
        val accuracy: 0.8575          val loss: 0.6535
Epoch 3      train accuracy: 0.9953  train loss: 0.0190

[I 2024-05-03 18:43:28,814] Trial 4 finished with value: 0.36414683094391453 and
parameters: {'hidden_dim': 95, 'cell_dropout': 0.09721499092775143}. Best is
trial 3 with value: 0.2906067268206523.

```

```

        val accuracy: 0.8600          val loss: 0.6309
Epoch 4      train accuracy: 0.9988  train loss: 0.0060
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
        val accuracy: 0.8500          val loss: 0.3127
Epoch 1      train accuracy: 0.6866  train loss: 0.5487
        val accuracy: 0.8850          val loss: 0.3032
Epoch 2      train accuracy: 0.9603  train loss: 0.1237
        val accuracy: 0.8900          val loss: 0.4365
Epoch 3      train accuracy: 0.9972  train loss: 0.0101
        val accuracy: 0.8875          val loss: 0.5352
Epoch 4      train accuracy: 1.0000  train loss: 0.0004

[I 2024-05-03 18:45:46,833] Trial 5 finished with value: 0.3031991250239886 and
parameters: {'hidden_dim': 94, 'cell_dropout': 0.21773799554081363}. Best is

```

trial 3 with value: 0.2906067268206523.

```
          val accuracy: 0.8950          val loss: 0.5613
Epoch 5      train accuracy: 1.0000  train loss: 0.0001
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
```

```
          val accuracy: 0.8000          val loss: 0.4602
Epoch 1      train accuracy: 0.6291  train loss: 0.6357
          val accuracy: 0.8575          val loss: 0.3929
Epoch 2      train accuracy: 0.9341  train loss: 0.1671
          val accuracy: 0.8725          val loss: 0.4704
Epoch 3      train accuracy: 0.9959  train loss: 0.0185
          val accuracy: 0.8525          val loss: 0.5169
Epoch 4      train accuracy: 0.9991  train loss: 0.0023
```

[I 2024-05-03 18:48:23,340] Trial 6 finished with value: 0.39287095846465003 and parameters: {'hidden_dim': 110, 'cell_dropout': 0.4919306658700668}. Best is trial 3 with value: 0.2906067268206523.

```
          val accuracy: 0.8550          val loss: 0.5912
Epoch 5      train accuracy: 0.9997  train loss: 0.0020
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
```

```
          val accuracy: 0.8625          val loss: 0.3172
Epoch 1      train accuracy: 0.6603  train loss: 0.5888
          val accuracy: 0.8625          val loss: 0.3585
Epoch 2      train accuracy: 0.9503  train loss: 0.1454
          val accuracy: 0.8900          val loss: 0.3525
Epoch 3      train accuracy: 0.9969  train loss: 0.0119
```

[I 2024-05-03 18:50:13,651] Trial 7 finished with value: 0.317195400595665 and parameters: {'hidden_dim': 86, 'cell_dropout': 0.48071266281454306}. Best is trial 3 with value: 0.2906067268206523.

```
          val accuracy: 0.8775          val loss: 0.6002
Epoch 4      train accuracy: 0.9988  train loss: 0.0053
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
```

```
          val accuracy: 0.8825          val loss: 0.3173
Epoch 1      train accuracy: 0.6775  train loss: 0.5740
          val accuracy: 0.8625          val loss: 0.4030
Epoch 2      train accuracy: 0.9491  train loss: 0.1461
          val accuracy: 0.8575          val loss: 0.5100
Epoch 3      train accuracy: 0.9925  train loss: 0.0243
```

[I 2024-05-03 18:52:25,688] Trial 8 finished with value: 0.3173267187980505 and parameters: {'hidden_dim': 119, 'cell_dropout': 0.12782325868286398}. Best is trial 3 with value: 0.2906067268206523.

```

                val accuracy: 0.8850          val loss: 0.5306
Epoch 4         train accuracy: 0.9994  train loss: 0.0024
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
                val accuracy: 0.8475          val loss: 0.3492
Epoch 1         train accuracy: 0.6381  train loss: 0.6063
                val accuracy: 0.8825          val loss: 0.3952
Epoch 2         train accuracy: 0.9403  train loss: 0.1601
                val accuracy: 0.8625          val loss: 0.6310
Epoch 3         train accuracy: 0.9953  train loss: 0.0130

[I 2024-05-03 18:53:54,522] Trial 9 finished with value: 0.34920736918082607 and
parameters: {'hidden_dim': 85, 'cell_dropout': 0.4157959216767957}. Best is
trial 3 with value: 0.2906067268206523.

```

```

                val accuracy: 0.8750          val loss: 0.6275
Epoch 4         train accuracy: 0.9997  train loss: 0.0017
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
                val accuracy: 0.8650          val loss: 0.3193
Epoch 1         train accuracy: 0.6916  train loss: 0.5649
                val accuracy: 0.8950          val loss: 0.2687
Epoch 2         train accuracy: 0.9550  train loss: 0.1313
                val accuracy: 0.8800          val loss: 0.4207
Epoch 3         train accuracy: 0.9969  train loss: 0.0104
                val accuracy: 0.8850          val loss: 0.5276
Epoch 4         train accuracy: 1.0000  train loss: 0.0005

[I 2024-05-03 18:55:16,659] Trial 10 finished with value: 0.2687366845516058 and
parameters: {'hidden_dim': 64, 'cell_dropout': 0.004205943985287786}. Best is
trial 10 with value: 0.2687366845516058.

```

```

                val accuracy: 0.8850          val loss: 0.5521
Epoch 5         train accuracy: 1.0000  train loss: 0.0001
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
                val accuracy: 0.8725          val loss: 0.3056
Epoch 1         train accuracy: 0.6853  train loss: 0.5735
                val accuracy: 0.8525          val loss: 0.4082
Epoch 2         train accuracy: 0.9509  train loss: 0.1264
                val accuracy: 0.8700          val loss: 0.5304
Epoch 3         train accuracy: 0.9969  train loss: 0.0092

[I 2024-05-03 18:56:21,811] Trial 11 finished with value: 0.30563747653594386
and parameters: {'hidden_dim': 64, 'cell_dropout': 0.0007470956028043929}. Best
is trial 10 with value: 0.2687366845516058.

```

```

                val accuracy: 0.8900          val loss: 0.4971

```

Epoch 4 train accuracy: 0.9997 train loss: 0.0008
 Early stopping triggered after 3 epochs with no validation loss improvement.
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

	val accuracy: 0.8125	val loss: 0.4242
--	----------------------	------------------

Epoch 1 train accuracy: 0.6681 train loss: 0.5680

	val accuracy: 0.8825	val loss: 0.3073
--	----------------------	------------------

Epoch 2 train accuracy: 0.9522 train loss: 0.1203

	val accuracy: 0.8900	val loss: 0.3991
--	----------------------	------------------

Epoch 3 train accuracy: 0.9975 train loss: 0.0066

	val accuracy: 0.8875	val loss: 0.4772
--	----------------------	------------------

Epoch 4 train accuracy: 1.0000 train loss: 0.0003

[I 2024-05-03 18:57:48,143] Trial 12 finished with value: 0.30734442747556245 and parameters: {'hidden_dim': 65, 'cell_dropout': 0.2873650476734631}. Best is trial 10 with value: 0.2687366845516058.

	val accuracy: 0.8875	val loss: 0.4814
--	----------------------	------------------

Epoch 5 train accuracy: 1.0000 train loss: 0.0001
 Early stopping triggered after 3 epochs with no validation loss improvement.
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

	val accuracy: 0.8625	val loss: 0.3541
--	----------------------	------------------

Epoch 1 train accuracy: 0.6309 train loss: 0.6173

	val accuracy: 0.8725	val loss: 0.3388
--	----------------------	------------------

Epoch 2 train accuracy: 0.9350 train loss: 0.1659

	val accuracy: 0.8875	val loss: 0.3496
--	----------------------	------------------

Epoch 3 train accuracy: 0.9959 train loss: 0.0172

	val accuracy: 0.9075	val loss: 0.4217
--	----------------------	------------------

Epoch 4 train accuracy: 0.9997 train loss: 0.0015

[I 2024-05-03 18:59:29,910] Trial 13 finished with value: 0.3388075392979842 and parameters: {'hidden_dim': 74, 'cell_dropout': 0.02701841868838295}. Best is trial 10 with value: 0.2687366845516058.

	val accuracy: 0.8725	val loss: 0.6391
--	----------------------	------------------

Epoch 5 train accuracy: 0.9997 train loss: 0.0004
 Early stopping triggered after 3 epochs with no validation loss improvement.
 Initializing embedding layer with pretrained word embeddings...
 Initialized 29841/32363 word embeddings

	val accuracy: 0.8725	val loss: 0.3334
--	----------------------	------------------

Epoch 1 train accuracy: 0.6550 train loss: 0.6134

	val accuracy: 0.8700	val loss: 0.3543
--	----------------------	------------------

Epoch 2 train accuracy: 0.9506 train loss: 0.1288

	val accuracy: 0.8775	val loss: 0.4464
--	----------------------	------------------

Epoch 3 train accuracy: 0.9988 train loss: 0.0065

[I 2024-05-03 19:00:42,762] Trial 14 finished with value: 0.3333545235487131 and parameters: {'hidden_dim': 76, 'cell_dropout': 0.23069484476552762}. Best is trial 10 with value: 0.2687366845516058.

```

            val accuracy: 0.8850          val loss: 0.4821
Epoch 4      train accuracy: 0.9994  train loss: 0.0013
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
            val accuracy: 0.8350          val loss: 0.3626
Epoch 1      train accuracy: 0.6684  train loss: 0.5816
            val accuracy: 0.8750          val loss: 0.3196
Epoch 2      train accuracy: 0.9469  train loss: 0.1469
            val accuracy: 0.8550          val loss: 0.5691
Epoch 3      train accuracy: 0.9962  train loss: 0.0137
            val accuracy: 0.8775          val loss: 0.5871
Epoch 4      train accuracy: 0.9994  train loss: 0.0018

[I 2024-05-03 19:02:18,873] Trial 15 finished with value: 0.319586476454368 and
parameters: {'hidden_dim': 75, 'cell_dropout': 0.29610041223238853}. Best is
trial 10 with value: 0.2687366845516058.

```

```

            val accuracy: 0.8725          val loss: 0.6029
Epoch 5      train accuracy: 1.0000  train loss: 0.0001
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
            val accuracy: 0.8875          val loss: 0.2882
Epoch 1      train accuracy: 0.6803  train loss: 0.5710
            val accuracy: 0.8800          val loss: 0.3472
Epoch 2      train accuracy: 0.9431  train loss: 0.1455
            val accuracy: 0.8650          val loss: 0.4676
Epoch 3      train accuracy: 0.9978  train loss: 0.0101

[I 2024-05-03 19:03:47,237] Trial 16 finished with value: 0.28817003736129176
and parameters: {'hidden_dim': 85, 'cell_dropout': 0.07497566875230444}. Best is
trial 10 with value: 0.2687366845516058.

```

```

            val accuracy: 0.8725          val loss: 0.6357
Epoch 4      train accuracy: 0.9988  train loss: 0.0026
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
            val accuracy: 0.8750          val loss: 0.2831
Epoch 1      train accuracy: 0.6719  train loss: 0.5681
            val accuracy: 0.8775          val loss: 0.3147
Epoch 2      train accuracy: 0.9519  train loss: 0.1321
            val accuracy: 0.8625          val loss: 0.4742
Epoch 3      train accuracy: 0.9981  train loss: 0.0075

[I 2024-05-03 19:05:08,547] Trial 17 finished with value: 0.2830975456879689 and
parameters: {'hidden_dim': 70, 'cell_dropout': 0.07315598741137136}. Best is
trial 10 with value: 0.2687366845516058.

```

```

            val accuracy: 0.8900          val loss: 0.5268

```

```

Epoch 4          train accuracy: 1.0000  train loss: 0.0005
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
          val accuracy: 0.8500          val loss: 0.3432
Epoch 1          train accuracy: 0.6766  train loss: 0.5663
          val accuracy: 0.8925          val loss: 0.3098
Epoch 2          train accuracy: 0.9497  train loss: 0.1328
          val accuracy: 0.8800          val loss: 0.4782
Epoch 3          train accuracy: 0.9959  train loss: 0.0125
          val accuracy: 0.8800          val loss: 0.6038
Epoch 4          train accuracy: 0.9997  train loss: 0.0013

[I 2024-05-03 19:06:45,945] Trial 18 finished with value: 0.30981439695908475
and parameters: {'hidden_dim': 69, 'cell_dropout': 0.04854634791931167}. Best is
trial 10 with value: 0.2687366845516058.

```

```

          val accuracy: 0.8775          val loss: 0.6515
Epoch 5          train accuracy: 1.0000  train loss: 0.0001
Early stopping triggered after 3 epochs with no validation loss improvement.
Initializing embedding layer with pretrained word embeddings...
Initialized 29841/32363 word embeddings
          val accuracy: 0.8200          val loss: 0.4029
Epoch 1          train accuracy: 0.6778  train loss: 0.5658
          val accuracy: 0.8900          val loss: 0.3172
Epoch 2          train accuracy: 0.9487  train loss: 0.1415
          val accuracy: 0.8600          val loss: 0.4811
Epoch 3          train accuracy: 0.9969  train loss: 0.0128
          val accuracy: 0.8900          val loss: 0.5496
Epoch 4          train accuracy: 0.9991  train loss: 0.0036

[I 2024-05-03 19:08:24,215] Trial 19 finished with value: 0.3171897759804359 and
parameters: {'hidden_dim': 69, 'cell_dropout': 0.131424981943556}. Best is trial
10 with value: 0.2687366845516058.

```

```

          val accuracy: 0.8850          val loss: 0.6746
Epoch 5          train accuracy: 1.0000  train loss: 0.0002
Early stopping triggered after 3 epochs with no validation loss improvement.

```

```

[58]: def reviews_train(
    model,
    train_loader,
    val_loader,
    lr=0.01,
    epochs=3,
    max_norm=5,
    loss_fn=torch.nn.functional.binary_cross_entropy,
    patience = 3,
):

```

```

model = model.to(DEVICE)

optimizer = torch.optim.Adam(
    [param for param in model.parameters() if param.requires_grad], lr=lr
)

best_val_loss = float("inf")
patience_counter = 0
for epoch in range(epochs):
    total_correct = total_loss = 0
    for reviews, labels in train_loader:
        model.train()

        reviews, labels = reviews.to(DEVICE), labels.to(DEVICE)

        prediction, _ = model(reviews)
        prediction = prediction.view(-1)
        loss = loss_fn(prediction, labels.float())

        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm =
↪max_norm)
        optimizer.step()
        optimizer.zero_grad()

        total_loss = total_loss + loss.item()
        prediction_class = torch.round(prediction)
        correct = (prediction_class == labels).sum().item()
        total_correct = total_correct + correct

    val_loss = reviews_eval(model, val_loader, loss_fn=loss_fn)

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        torch.save(model.state_dict(), f"model_best_val_loss.pt")
        patience_counter = 0
    else:
        patience_counter += 1
    print(
        f"Epoch {epoch + 1:2}\t"
        f"train accuracy: {total_correct / len(train_loader.dataset):.4f}\t"
        f"train loss: {total_loss / len(train_loader):.4f}"
    )
    if patience_counter >= patience:
        print(f"Early stopping triggered after {patience} epochs with no
↪validation loss improvement.")
        break

```



```
reviews_eval(model, val_loader, loss_fn=loss_fn)
```

```
[59]: vocab_size = len(dataset.vocab)
      embedding_dim = 100
      hidden_dim = 64
      num_layers = 2
      n_epochs = 10
      cell_type = "gru"
      bidirectional = True
      cell_dropout = 0.004205943985287786
```

```
[64]: model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
      ↪cell_type, cell_dropout, bidirectional)
      reviews_load_embeddings(model_pf.embedding, vocab.get_stoi())
      reviews_train(model_pf, train_loader, val_loader, epochs=n_epochs)
```

Initializing embedding layer with pretrained word embeddings...

Initialized 29841/32363 word embeddings


	val accuracy: 0.8675	val loss: 0.3289
Epoch 1	train accuracy: 0.6916	train loss: 0.5717
	val accuracy: 0.8675	val loss: 0.3289
	val accuracy: 0.8950	val loss: 0.2899
Epoch 2	train accuracy: 0.9491	train loss: 0.1452
	val accuracy: 0.8950	val loss: 0.2899
	val accuracy: 0.8725	val loss: 0.5185
Epoch 3	train accuracy: 0.9988	train loss: 0.0059
	val accuracy: 0.8725	val loss: 0.5185
	val accuracy: 0.8950	val loss: 0.4998
Epoch 4	train accuracy: 0.9997	train loss: 0.0008
	val accuracy: 0.8950	val loss: 0.4998
	val accuracy: 0.8950	val loss: 0.5114
Epoch 5	train accuracy: 1.0000	train loss: 0.0001

Early stopping triggered after 3 epochs with no validation loss improvement.


```
[65]: model_pf = SimpleLSTM(vocab_size, embedding_dim, hidden_dim, num_layers,
      ↪cell_type, cell_dropout, bidirectional)
      model_pf.load_state_dict(torch.load("model_best_val_loss.pt"))
      print(" testing loss: ", reviews_eval(model_pf, test_loader, loss_fn = torch.nn.
      ↪functional.binary_cross_entropy))
```

```
val accuracy: 0.8650      val loss: 0.3505
testing loss: 0.3504867765765924
```

```
[66]: nextplot()
      _ = tsne_thought(model_pf, train_loader, DEVICE)
```



```
[67]: nextplot()  
_ = tsne_thought(model_pf, val_loader, DEVICE)
```



```
[68]: nextplot()  
_ = tsne_thought(model_pf, test_loader, DEVICE)
```

