

Deep Learning Assignment 2

Kok Teng, Ng (1936360), nkokteng


MinJeong, Lee (1978925), minjlee

Question 4: Pretrained Embedding & Visualization

- a. The file **“data/word-embedding.txt”** contains pre-trained word embeddings as known as GloVe, which are essentially numerical representations of words in a high-dimensional space. These embeddings are generated for a vocabulary set of 29,841 words, each *represented* as a **dense, 100-dimensional vector**. These **word embeddings** are the *product of* a specific type of neural network as **embedding neural network**. The primary function of this network is to transform sparse, high-dimensional word vector (*usually one-hot encoding vectors*) into a lower-dimensional, **dense format**. This transformation is achieved through the learning process (such as **Skip-Gram model or Continuous Bag-of-Words (CBOW)**), where the network adjusts the word vectors based on the contextual usage of words within a large corpus of text. The primary advantage of these dense vectors is *their capability to capture semantic relationships between words*. This is achieved using **cosine similarity**, a metric that *measure the cosine of the angle between two vectors*. In the context of word embeddings, a **high cosine similarity score** between two-word vectors *indicates* that the words are used in *similar contexts and having same direction and close to each other within the embedding (latent representation) space*, suggesting a close semantic relationship. Conversely, dissimilar words have vectors that diverge, resulting in a negative similarity score.

One of the most fascinating aspects of these word embedding is their ability to **capture complex semantic relationship** through **simple vector arithmetic**. For instance, these embedding’s vector arithmetic operation: **“king” – “man” + “woman” = “queen”**. This operation demonstrates that the relationships between words, such as **gender**, can be represented and manipulated within the embedding space.

- b. Based on the bottom illustration, the t-Distribution Stochastic Neighbor Embedding (t-SNE) visualization is a technique used to visualize high-dimensional data (in our case, the high-dimensional data is regarding to the word representation vector) in a two- or three-dimensional map. This method allows us to inspect the structure of high-dimensional data more intuitively.



In the context of word representation vectors, tokens with similar semantics in the high-dimensional space, the goal of t-SNE will ensure and aim to map them closer together

in the lower-dimensional visualization. Investigate from the illustration, we could view that the pair of “*film*” and “*movie*”, “*i*”, “*me*”, and “*my*”, “*a*” and “*an*”, and others are having the same semantic, and they were also being mapped close to each other in the lower-dimensional by using the method of t-SNE to plot them in the 2-Dimensional visualization. Additionally, the tokens of “*they*” and “*i*” are having different semantic since the semantics of “*they*” is regarding to many people and “*i*” is regarding to one people; therefore, both tokens are having different direction to each other as getting a large distance.

The measure of similarity between tokens is determined using an isotropic Gaussian distribution, which essentially investigates the similarity of tokens in the high-dimensional (embedding) space. The isotropic Gaussian $P(z_j | z_i) = \mathcal{N}(z_j | z_i, \sigma^2 I)$. Based on this Gaussian formula, we could plot the distribution of a given token “*j*” and accesses the density of other tokens within this distribution. If another token “*i*” lies within a high-density region of the distribution of token “*j*”, it is close to “*j*”.




During the computation of similarity, the variance of the Gaussian distribution plays a crucial role. It controls the spread of the distribution, thereby determining which tokens are considered neighbors or close to the target token “*j*”. A small variance of Gaussian distribution results in a small group of data points being labeled as having similar embedding values or being close to each other. Conversely, a large variance might include a larger group of data points as having similar embedding values, but there might also include noise into the data. To determine the variance, needed to compute the Shannon Entropy H_i of the neighbor distribution $P(z_j | z_i)$. Intuitively, this represents the number of effective neighbors of data point “*i*” in bits. Compute the perplexity of the distribution from the Shannon Entropy as 2^{H_i} . Perplexity could be thought of as a measure of how well a probability distribution predicts a sample. Use the target perplexity to determine the variance of the *t-SNE*’s *isotropic Gaussian*.

Additionally, we are utilizing the **Uniform Manifold Approximation and Projection (UMAP)** to reduce the dimensionality of word embeddings for visualization purposes. UMAP is favored for its consistency across different runs, unlike t-SNE which can produce different visualizations with the same data due to its gradient descent procedure that can sometimes get stuck in local minima.

Due to **UMAP** uses spectral embedding to compute the probability of neighboring data points, making it more stable and less likely to be trapped in local minima. This stability is crucial ensuring the reliability of our visualizations.

In our exploration, we used **UMAP** to project word embeddings into 2-Dimensional space. We observed that semantically similar tokens, such as “*positively*” and “*appreciate*”, “*reproduction*” and “*duplicated*”, and others were mapped close to each other in the **UMAP** visualization. This demonstrates **UMAP**’s effectiveness in

preserving semantic relationships when reducing the dimensionality of word embeddings.




- c. During the training with the specified hyperparameters on the *SimpleRNN* model for 10 epochs, we observed the potential **overfitting** issue. This was indicated from the low loss on the training set and the comparatively high loss on the validation set. Therefore, we can infer that the model's performance on the validation set might be poor. This could also reflect in the t-SNE visualization by using the thought vector of each review from validation set, which appear more dispersed for the validation set compared to the training set according to the label of the review. Additionally, we could also investigate from the result as the model was starting to overfitting after the epoch 3 as the loss from the validation set was starting to increase but the loss from training set was still decreasing.

Epoch 1	train accuracy: 0.4928	train loss: 0.6999	Epoch 6	train accuracy: 0.6508	train loss: 0.5464
	val accuracy: 0.4825	val loss: 0.6957		val accuracy: 0.4959	val loss: 0.8098
Epoch 2	train accuracy: 0.4928	train loss: 0.6938	Epoch 7	train accuracy: 0.7472	train loss: 0.4570
	val accuracy: 0.4825	val loss: 0.6934		val accuracy: 0.6175	val loss: 0.8547
Epoch 3	train accuracy: 0.5186	train loss: 0.6929	Epoch 8	train accuracy: 0.8853	train loss: 0.2865
	val accuracy: 0.4725	val loss: 0.6939		val accuracy: 0.6675	val loss: 0.8698
Epoch 4	train accuracy: 0.5316	train loss: 0.6897	Epoch 9	train accuracy: 0.9580	train loss: 0.1435
	val accuracy: 0.5080	val loss: 0.6998		val accuracy: 0.7125	val loss: 0.8564
Epoch 5	train accuracy: 0.6089	train loss: 0.6364	Epoch 10	train accuracy: 0.9812	train loss: 0.0540
	val accuracy: 0.5208	val loss: 0.7398		val accuracy: 0.7359	val loss: 0.9660

In the context of *t-SNE* visualization (by using *tsne_thought()* function), we will only use the last hidden states (also known as “*thought vectors*”, which is the vector of prediction with respect to the reviews before fitting into the logit function) of the trained model, and project into a lower-dimensional space as 2-dimensional.

During the feedforward process, reviews sharing the same label are expected to yield similar or close thought vectors. This is because these vectors, which represent the final hidden state, compress the information from all previous time steps in the sequence.

When these thought vectors are passed through the logit function, they are likely to produce the same label. Therefore, in the lower-dimensional space created by *t-SNE*, reviews with similar content (will also having the similar thought vector) are expected to be mapped closer together in the lower-dimensional. This visualization could provide insight how the model is internally processing and categorizing the reviews.



(a) training set

(b) validation set

Based on the above t-SNE visualization, it appears that the trained model is performing as expected. Reviews with the same label tend to have similar “thought vectors” produced by the feedforward RNN sequence model. When these vectors are projected into a lower-dimensional and visualized using t-SNE, reviews with the same labels are mapped closer to each other in the lower-dimensional in the training set. This suggests that the model is effectively compressing and summarizing the information from all previous time steps in the sequence.

However, we could also investigate the model’s performance is having the signal of **overfitting**, as the model doesn’t perform as well on the validation set. This could be due to the model doesn’t effectively capturing and compressing the information from all previous time steps in the sequence to generate the thought vectors. Therefore, reviews with the same label do not always generate similar or close vector representation (also called “*thought vector*”) and predict it wrongly. This results in a sparse distribution of data points in the *t-SNE* visualization, indicating that the model struggles to map reviews with similar labels closer to each other within the lower-dimensional due to there are having dissimilar thought vector with the same label of review. Therefore, the model can’t promise that it could generalize well on those unseen data.

- d. In this question, the pretrained *GloVe embeddings model* is incorporated as the *embedding layer* within the custom *SimpleLSTM sequence model*. This integration enables the model to transform categorical input data, represented as *One-Hot Encoding*, into *meaningful word embeddings*.

Upon loading the pretrained embeddings, the model is subsequently trained on a specific task – in this case, sentiment analysis. During this training phase, the weights of *the embedding layer are fine-tuned*. This fine-tuning process is accomplished by computing the *gradient of the embedding layer’s weights with respect to the loss function*. These gradients are then used to *update and adjust the weights*, enabling them to better align with the specific task at hand. This fine-tuning process allows the model to leverage the general language understanding capabilities of the pretrained embeddings while also adapting to the nuances of the specific task. It is often observed that fine-tuning the embedding layer leads to *better performance compared to the* upcoming case where the embedding layer is *not fine-tuned*.

Epoch 1	train accuracy: 0.4888	train loss: 0.6950	Epoch 6	train accuracy: 0.8213	train loss: 0.3973
	val accuracy: 0.4858	val loss: 0.6932		val accuracy: 0.7556	val loss: 0.5638
Epoch 2	train accuracy: 0.5153	train loss: 0.6985	Epoch 7	train accuracy: 0.9631	train loss: 0.2668
	val accuracy: 0.5288	val loss: 0.6928		val accuracy: 0.7825	val loss: 0.5683
Epoch 3	train accuracy: 0.5513	train loss: 0.6928	Epoch 8	train accuracy: 0.9337	train loss: 0.3984
	val accuracy: 0.4828	val loss: 0.6932		val accuracy: 0.7958	val loss: 0.6733
Epoch 4	train accuracy: 0.5425	train loss: 0.6826	Epoch 9	train accuracy: 0.9547	train loss: 0.1324
	val accuracy: 0.4958	val loss: 0.6943		val accuracy: 0.7458	val loss: 0.7889
Epoch 5	train accuracy: 0.6425	train loss: 0.5942	Epoch 10	train accuracy: 0.9716	train loss: 0.0886
	val accuracy: 0.6989	val loss: 0.6207		val accuracy: 0.8888	val loss: 0.1236

The results presented indicate that the model is trained effectively on the training set, as evidenced by the low loss and high accuracy. However, it appears that the model begins to **overfit** after the fourth epoch. This is suggested by the increasing validation loss, despite the model's continued high performance on the training set. This indicates that the model is becoming excessively specialized to the training data, thereby losing its capability to generalize effectively on unseen data.


- e. In this question, the **pretrained GloVe embeddings model** is once again employed as **embedding layer** within the **custom SimpleLSTM sequence model**. However, in contrast to the previous case, the parameters of this embedding layer are **not subject to fine-tuning**. This implies that the model continues to **leverage the general language understanding** derived from the pretrained embeddings to **perform the current task**.

Epoch 1	train accuracy: 0.4928	train loss: 0.6953	Epoch 6	train accuracy: 0.5766	train loss: 0.6513
	val accuracy: 0.5125	val loss: 0.6931		val accuracy: 0.5825	val loss: 0.7334
Epoch 2	train accuracy: 0.5141	train loss: 0.6929	Epoch 7	train accuracy: 0.6875	train loss: 0.6163
	val accuracy: 0.5288	val loss: 0.6919		val accuracy: 0.5825	val loss: 0.7670
Epoch 3	train accuracy: 0.5162	train loss: 0.6896	Epoch 8	train accuracy: 0.6972	train loss: 0.5623
	val accuracy: 0.4875	val loss: 0.8386		val accuracy: 0.7125	val loss: 0.6833
Epoch 4	train accuracy: 0.5262	train loss: 0.6957	Epoch 9	train accuracy: 0.7389	train loss: 0.5567
	val accuracy: 0.4988	val loss: 0.6936		val accuracy: 0.6875	val loss: 0.6195
Epoch 5	train accuracy: 0.5509	train loss: 0.6783	Epoch 10	train accuracy: 0.7583	train loss: 0.5176
	val accuracy: 0.5188	val loss: 0.6943		val accuracy: 0.6875	val loss: 0.6246

Based on the results presented, it can be inferred that the model may be **underfitting**. This is because the parameters of the embedding layer are not adjusted during the training process to better suit the specific task. Therefore, the model relies solely on the general language understanding to execute the task. During the backward propagation phase, the model computes the gradient of the embedding layer's parameter with respect to the loss function and updates these parameters to minimize the loss generated by the model during forward pass. However, without fine-tuning these parameters, the model may struggle to generalize effectively to unseen data and training data, leading to potential underfitting problem.

Question 5: Exploration

- a. There are having seven different setups based on the **SimpleLSTM model** as blueprint.




Bidirectional LSTM with 0.0 dropout

The model, which is configured with a **Bidirectional LSTM cell** and a **0,0-dropout rate** begins to exhibit signs of **overfitting after the fourth epoch**. This is evidenced by the model's superior performance on the training set, contrasted with its inability to generalize effectively on the validation set.

The t-SNE visualization of the model's thought vector further substantiates this observation. In the training set, the model successfully generates closely mapped representations for reviews with the same label in the lower-dimensional space. However, it struggles to consistently cluster reviews with the same label in proximity,

indicating a lack of generalization. This issue becomes more pronounced in the validation set, where the model fails to map similar reviews closely together.

This discrepancy between the training and validation sets is clear indication of the model's tendency to overfit. In other words, while the model is highly adept at learning the specific characteristic of the training data, it struggles to apply this knowledge effectively to unseen data, thereby compromising its predictive accuracy.



Bidirectional LSTM with 0.8 dropout

The model, configured with a **Bidirectional LSTM cell** and a **0,8-dropout rate**. This dropout technique, which randomly deactivates a subset of neurons during each training epoch, effectively reduces the model's complexity by preventing certain neurons from updating their parameters. This strategy aids in mitigating overfitting.

However, the implementation of dropout, signs of overfitting are still observed in this model, but it has **less overfitting compared to the previous case** as the validation set's loss are lower than it. This is evident in the t-SNE visualization of the training set, where the model adeptly generates similar thought vectors for reviews with the same label, mapping them closely in the lower-dimensional space. This indicates that the model has learned the training set well. However, it struggles to consistently cluster reviews within the same label closely together. When we examine the t-SNE visualization of the validation set, it becomes clear that the model struggles slightly to generalize its learning to unseen data. Although the model's performance on the validation set has **improved compared to the model without dropout**, there are still instances where data points with different labels are mapped closely together in the lower-dimensional space. This suggests that while the **dropout technique has enhanced** the model's capability to **generalize**.

We also compare the performance of model configured with **Unidirectional and Bidirectional LSTM cell**, each either with **0,0- or 0,8-dropout rate**. The results indicate that the **Bidirectional LSTM models outperform their Unidirectional counterparts**, irrespective of the dropout rate. This **performance difference** can be attributed to the inherent **characteristics of LSTM cells**.


Unidirectional LSTM models process information in a **single direction** along the input sequence data. However, in the context of our task, which involves processing review texts, **both preceding and succeeding tokens can influence the semantics of a sentence**. Unidirectional LSTM model is **unable to capture** this bidirectional context, which may **limit their performance**.

In contrast, **Bidirectional LSTM models** process the sentence in **both directions**, thereby capturing a **more comprehensive understanding** of the sequence. The


bidirectional processing capability is likely *why Bidirectional LSTM models outperform Unidirectional LSTM model.*

Epoch 1	train accuracy: 0.9975	train loss: 0.4915
	val accuracy: 0.4825	val loss: 0.4934
Epoch 2	train accuracy: 0.9987	train loss: 0.4769
	val accuracy: 0.5006	val loss: 0.4963
Epoch 3	train accuracy: 0.9967	train loss: 0.4628
	val accuracy: 0.7256	val loss: 0.3789
Epoch 4	train accuracy: 0.9925	train loss: 0.3759
	val accuracy: 0.8638	val loss: 0.3623
Epoch 5	train accuracy: 0.9828	train loss: 0.3614
	val accuracy: 0.7986	val loss: 0.4953
Epoch 6	train accuracy: 0.9944	train loss: 0.3234
	val accuracy: 0.8386	val loss: 0.7986
Epoch 7	train accuracy: 0.9958	train loss: 0.3312
	val accuracy: 0.8633	val loss: 0.4613
Epoch 8	train accuracy: 0.9966	train loss: 0.3664
	val accuracy: 0.8275	val loss: 0.4840
Epoch 9	train accuracy: 0.9994	train loss: 0.3637
	val accuracy: 0.8625	val loss: 1.5207
Epoch 10	train accuracy: 0.9977	train loss: 0.3628
	val accuracy: 0.8108	val loss: 1.1340

Performance



t-SNE training set




t-SNE validation set

Unidirectional LSTM with 0.0 dropout


The *Unidirectional LSTM model with a 0,0-dropout rate* demonstrates good generalization on the training set, as evidenced by the t-SNE visualization. Reviews with the same label are mapped closely in the lower-dimensional space, indicating that the model has learned to generate similar thought vectors for these reviews. However, on the validation set, different labels are mapped closely together, suggestion that the model struggles to generalize well to unseen data.

Epoch 1	train accuracy: 0.9996	train loss: 0.4987
	val accuracy: 0.5476	val loss: 0.4958
Epoch 2	train accuracy: 0.9985	train loss: 0.4889
	val accuracy: 0.8068	val loss: 0.4951
Epoch 3	train accuracy: 0.9917	train loss: 0.4834
	val accuracy: 0.9779	val loss: 0.4948
Epoch 4	train accuracy: 0.9949	train loss: 0.4711
	val accuracy: 0.8388	val loss: 0.7221
Epoch 5	train accuracy: 0.9942	train loss: 0.3129
	val accuracy: 0.7176	val loss: 0.5057
Epoch 6	train accuracy: 0.9854	train loss: 0.3847
	val accuracy: 0.7979	val loss: 0.4934
Epoch 7	train accuracy: 0.9939	train loss: 0.2774
	val accuracy: 0.7957	val loss: 0.4973
Epoch 8	train accuracy: 0.9964	train loss: 0.3244
	val accuracy: 0.7970	val loss: 0.7533
Epoch 9	train accuracy: 0.9945	train loss: 0.3389
	val accuracy: 0.7926	val loss: 0.7643
Epoch 10	train accuracy: 0.9966	train loss: 0.3159
	val accuracy: 0.7902	val loss: 0.4733

Performance



t-SNE training set




t-SNE validation set

Unidirectional LSTM with 0.8 dropout


The *Unidirectional LSTM model with a 0,8-dropout rate* appears to *slightly underfitting*, as indicated by t-SNE visualization of the training set. The model struggles to generate similar thought vectors for reviews with the same label, resulting in mixed classes within the same cluster in the lower-dimensional space. This issue could be due to high dropout rate, which deactivates some neurons during training, potentially preventing the model from learning necessary patterns from the data. This underfitting is further confirmed by the model's performance on the validation set, where it fails to accurately map reviews with the same labels.

Epoch 1	train accuracy: 0.9703	train loss: 0.4977
	val accuracy: 0.8088	val loss: 0.4174
Epoch 2	train accuracy: 0.9128	train loss: 0.2386
	val accuracy: 0.8075	val loss: 0.3963
Epoch 3	train accuracy: 0.9638	train loss: 0.3076
	val accuracy: 0.8688	val loss: 0.6587
Epoch 4	train accuracy: 0.9764	train loss: 0.4880
	val accuracy: 0.8688	val loss: 0.4123
Epoch 5	train accuracy: 1.0000	train loss: 0.4680
	val accuracy: 0.8625	val loss: 0.7137
Epoch 6	train accuracy: 1.0000	train loss: 0.3953
	val accuracy: 0.8688	val loss: 0.2796
Epoch 7	train accuracy: 1.0000	train loss: 0.3864
	val accuracy: 0.8688	val loss: 0.4897
Epoch 8	train accuracy: 1.0000	train loss: 0.3883
	val accuracy: 0.8623	val loss: 0.8348
Epoch 9	train accuracy: 1.0000	train loss: 0.3688
	val accuracy: 0.8656	val loss: 0.8688
Epoch 10	train accuracy: 1.0000	train loss: 0.3958
	val accuracy: 0.8620	val loss: 0.3855

Performance



t-SNE training set



t-SNE validation set

Unidirectional GRU with 0.0 dropout


The model, configured with a *Unidirectional GRU cell* and a *0,0-dropout rate*, demonstrates *superior performance on the training set compared to models utilizing either Bidirectional or Unidirectional LSTM cells*. This is substantiated by the t-SNE visualization, where the Unidirectional GRU model generates *thought vectors for*

reviews with the same label that are more closely clustered in the lower-dimensional space compared to those generated by the LSTM cell models. This suggests that the GRU model is more effective in generalizing the thought vectors for reviews with the same label within the context of the current task.


However, signs of overfitting begin to emerge after the second epoch too. This is indicated by a decreasing loss for training set, contrasted with an increasing loss for the validation set. The t-SNE visualization of the validation set further confirms this, as data points with the same label are slightly more dispersed.

Epoch 1	train accuracy: 0.9754	train loss: 0.4512
	val accuracy: 0.8188	val loss: 0.4138
Epoch 2	train accuracy: 0.9975	train loss: 0.2649
	val accuracy: 0.8625	val loss: 0.3903
Epoch 3	train accuracy: 0.9912	train loss: 0.4066
	val accuracy: 0.8525	val loss: 0.3956
Epoch 4	train accuracy: 0.9991	train loss: 0.4013
	val accuracy: 0.8958	val loss: 0.4438
Epoch 5	train accuracy: 0.9994	train loss: 0.4625
	val accuracy: 0.8758	val loss: 0.6099
Epoch 6	train accuracy: 0.9997	train loss: 0.4038
	val accuracy: 0.8575	val loss: 0.8438
Epoch 7	train accuracy: 1.0000	train loss: 0.4001
	val accuracy: 0.8650	val loss: 0.7408
Epoch 8	train accuracy: 1.0000	train loss: 0.4001
	val accuracy: 0.8650	val loss: 0.7716
Epoch 9	train accuracy: 1.0000	train loss: 0.4000
	val accuracy: 0.8650	val loss: 0.7972
Epoch 10	train accuracy: 1.0000	train loss: 0.4000
	val accuracy: 0.8650	val loss: 0.8168

Performance



t-SNE training set




t-SNE validation set

Bidirectional GRU with 0.8 dropout


The model, configured with a **Bidirectional GRU cell** and a **0,8-dropout rate**, demonstrates impressive performance on the training set. This is substantiated by the t-SNE visualization, where the model generates thought vectors for reviews with the same label that are closely clustered in the lower-dimensional space. This means that the model is effectively learning to generate similar thought vectors for reviews with the same label, thereby forming distinct clusters. However, signs of **overfitting** begin to emerge after the second epoch.

Epoch 1	train accuracy: 0.8108	train loss: 0.7033
	val accuracy: 0.4888	val loss: 0.7259
Epoch 2	train accuracy: 0.5878	train loss: 0.7013
	val accuracy: 0.5008	val loss: 0.6996
Epoch 3	train accuracy: 0.5162	train loss: 0.6981
	val accuracy: 0.4779	val loss: 0.7132
Epoch 4	train accuracy: 0.4733	train loss: 0.6861
	val accuracy: 0.5176	val loss: 0.7182
Epoch 5	train accuracy: 0.6132	train loss: 0.6356
	val accuracy: 0.4675	val loss: 0.7285
Epoch 6	train accuracy: 0.6416	train loss: 0.5761
	val accuracy: 0.5129	val loss: 0.7961
Epoch 7	train accuracy: 0.6486	train loss: 0.5361
	val accuracy: 0.5125	val loss: 0.7994
Epoch 8	train accuracy: 0.6667	train loss: 0.5069
	val accuracy: 0.4868	val loss: 0.9076
Epoch 9	train accuracy: 0.6719	train loss: 0.4889
	val accuracy: 0.4951	val loss: 0.9097
Epoch 10	train accuracy: 0.6992	train loss: 0.4704
	val accuracy: 0.4858	val loss: 0.9999

Performance



t-SNE training set



t-SNE validation set

Unidirectional Elman (RNN) with 0.0 dropout

The model, configured with a **Unidirectional Elman cell** and **0,0-dropout rate**, having the least performance, and indicating a case of underfitting model. This suggests that the model is too simplistic to effectively handle the complexity of the current task. This is evident from the accuracy and loss metrics on both the validation and training sets, with the accuracy less than 0,7 and loss exceeding 0,45. The t-SNE visualization of the training and validation sets further confirms this. Reviews with different labels have similar thought vectors and are closely mapped in the lower-dimensional space. Conversely, reviews with the same label have dissimilar thought vector and are mapped further apart.

In conclusion, our comprehensive analysis indicates that the model configured with a **Bidirectional GRU cell**, and a **0,8-dropout rate** emerges as the **most optimal solution** for the task at hand, **outperforming the other six model configurations**. This conclusion is primarily drawn from its superior performance during the training and

validation phase and the t-SNE visualization. The t-SNE visualization of the thought vector for this mode, for both the training and validation sets, provides further evidence of its effectiveness. The visualization reveals that reviews sharing the same label are closely mapped to each other in the lower-dimensional space. This clustering pattern is a clear indication of the model's capability to effectively generalize and distinguish between different classes. In contrast, other models with *different hyperparameters are having slightly bad performance*, further underscoring the effectiveness of the chosen hyperparameter setup for the Bidirectional GRU model. Therefore, we recommend the use of a ***Bidirectional GRU cell with 0,8-dropout rate*** for this current task.

- b. In our endeavor to optimize the performance of our ***Bidirectional GRU cell model***, we have employed the ***Optuna*** framework for hyperparameter tuning. This process involves conducting 20 trials to ***identify the optimal hyperparameters*** that ***yield the lowest validation loss***. The loss function used for this purpose is ***Binary Cross Entropy***.

```
hidden_dim = trial.suggest_int('hidden_dim', 64, 128)
cell_dropout = trial.suggest_float('cell_dropout', 0.0, 0.5)
```

To ***mitigate*** the risk of ***overfitting***, we have incorporated ***early stopping mechanism***. This is particularly crucial as our investigations revealed a tendency for models to overfit the training set after few epochs, we set the ***patient as 3*** if the validation loss ***doesn't decrease more than 3 times*** then we ***pause the training*** and ***take the model's stat (weights)*** as returned model to ***test on testing set***. The hyperparameters under consideration during this tuning process include the ***hidden dimension and dropout rate***. This comprehensive approach ensures effective optimization of our model's performance.

The ***optimal hyperparameters*** identified for our ***Bidirectional GRU cell sequence model*** is ***hidden layer dimension of 64 and 0,00421-dropout rate***. These settings yield the ***lowest validation loss***.


```
val accuracy: 0.8576   val loss: 0.3287
Epoch 1: train accuracy: 0.8552   train loss: 0.3737
val accuracy: 0.8676   val loss: 0.3289
val accuracy: 0.8698   val loss: 0.3099
Epoch 2: train accuracy: 0.8681   train loss: 0.3182
val accuracy: 0.8958   val loss: 0.2899
val accuracy: 0.8725   val loss: 0.3180
Epoch 3: train accuracy: 0.8708   train loss: 0.3079
val accuracy: 0.8725   val loss: 0.3180
val accuracy: 0.8958   val loss: 0.3098
Epoch 4: train accuracy: 0.8977   train loss: 0.3008
val accuracy: 0.8958   val loss: 0.3098
val accuracy: 0.8958   val loss: 0.3124
Epoch 5: train accuracy: 1.0000   train loss: 0.4803
Early stopping triggered after 5 epochs with no validation loss improvement
```

The model training process, as depicted above, demonstrates an interesting pattern. In the ***initial epoch***, the model achieves a ***validation loss of 0,328***, which further reduces to ***0,2899 in the second epoch***. However, ***post*** the second epoch, the model ***begins to overfit***, leading to ***an increase in validation loss***. This overfitting indicates that the model is not generalizing well to the validation set, resulting in a ***higher loss***.

To mitigate this issue, we employ an ***early abort mechanism***. This mechanism monitors the ***validation loss*** and ***halts*** the training process if the ***validation loss doesn't improve*** over a certain number of epochs. In our setting, the training is stopped if there is no improvement in validation loss for ***three consecutive epochs***.


The ***parameters*** of the model at the point of ***lowest validation loss*** are saved in a ***PyTorch file***. This allows us to later load the model with the optimal parameters, ensuring that we utilize the most effective version of the model for our predictions. This approach helps in preventing utilize the overfitting model to generalize on unseen data

and achieving a balance between bias and variance, thereby improving the model's performance on unseen data.



The above t-SNE visualization of the thought vectors generated from *the training set* using these *optimal hyperparameters and model's parameters with least validation loss within the training epoch* provides further insights. It is evident that the model can generate similar thought vector for reviews that share the same label. Conversely, dissimilar thought vectors are generated for reviews with different labels. However, there are few instances where similar thought vectors are generated for reviews with different labels. This is an expected outcome considering our use of early stopping based on validation loss. The model used for this visualization is the one that achieved the lowest validation loss, but the training accuracy and loss haven't reached their optimal score. While it may not be perfectly accurate, it strikes a balance between avoiding overfitting and underfitting, thereby ensuring a robust model performance.

val accuracy: 0.8650 val loss: 0.3505
testing loss: 0.3504867765765924



Testing set accuracy and validation and t-SNE visualization

When it comes to the testing set, the model achieves a **testing accuracy score of 0,8650** and **testing loss of 0,3505**. These results are *slightly lower than those achieved on the validation set*, which is expected as the model has not seen the testing data during training. The t-SNE visualization of the thought vectors generated from the testing set would likely show *similar pattern to the validation set*, with the model generating similar thought vectors for reviews that share the same label and dissimilar thought vectors for reviews with different labels. However, given the slightly lower performance on the testing set, we could observe a slightly higher degree of overlap between thought vectors of different labels from the t-SNE visualization.